

# Introduction To Operating System

## Operating System

An **Operating System (OS)** is the interface between the **user** and the **computer hardware**. It manages and coordinates the use of hardware attached to the computer system.

- **Goals of an Operating System:**

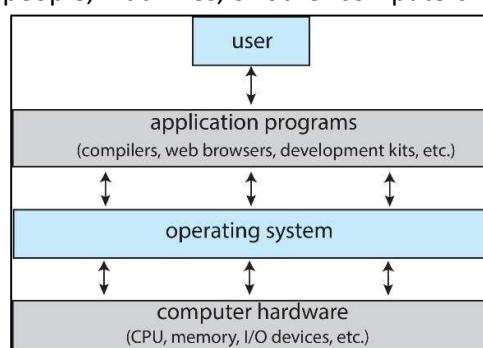
1. Execute user programs:
2. Make solving user problems easier:
3. Make the computer system convenient to use:
4. Use the computer hardware in an efficient manner

- **Tasks Performed by the OS:**

1. Memory Management:
2. Peripheral Device Control:
3. Handling Input and Output:
4. File Management:

## Components of a Computer System

1. **Hardware:** Basic computing resources: CPU, memory, and I/O devices.
2. **Operating System:** Controls and coordinates hardware usage among applications and users.
3. **Application Programs:** Use system resources to solve user problems (e.g., word processors, web browsers, video games).
4. **Users:** Can be people, machines, or other computers.



## **Basic functions of operating systems:**

### **1. Memory Management**

It manages the allocation of memory of system for different processes. It manages both the primary memory and secondary memory.

### **2. Processor Management**

It manages all the running processes in computer system. A process is simply a program that is run by a user on computer system.

### **3. Security Management**

It ensures the security of computer system from the various threats and viruses attacks. An operating system uses various techniques such as authentication, authorization, cryptography etc. for ensuring security of computer system.

### **4. Device Management**

This function of operating system is used to manage different devices that are connected with the computer system. An operating system interacts with hardware device through specified device drivers.

### **5. File Management**

An operating system manages the files and directories of computer system. A file can be defined as a collection of information or data that is stored in the memory of computer system. An operating system allows us to create, delete, save, edit files in a computer system.

## **Process Management**

- A process is a program in execution and represents a unit of work within the system.
- A program is passive, while a process is active and requires resources like CPU, memory, I/O devices, and files.
- Initialization data may be needed for a process to start.
- When a process terminates, the operating system reclaims any reusable resources.
- A single-threaded process has one program counter that specifies the next instruction to execute and runs instructions sequentially until completion.
- A multi-threaded process has one program counter per thread, allowing multiple tasks to run simultaneously.
- Typically, many processes (user and operating system) run concurrently on one or more CPUs.
- Concurrency is achieved by multiplexing the CPUs among the processes and threads.

## Memory Management

Memory management is a crucial function of an operating system that ensures programs can run effectively.

- **Purpose:** To execute a program, all or part of its instructions and the necessary data must be loaded into memory.
- **Resource Management:** The operating system decides what data and processes are in memory at any time, optimizing CPU usage and improving response times for users.

### Key Activities in Memory Management:

1. **Tracking Memory Usage:** The OS keeps track of which parts of memory are currently in use and by which processes.
2. **Moving Data:** It decides which processes or data need to be moved into or out of memory based on their requirements.
3. **Allocating Memory:** The OS allocates memory space to processes when needed and deallocates it when the processes are finished, ensuring efficient use of memory resources.

## Storage Management

- The operating system provides a uniform and logical view of information storage.
- It abstracts the physical properties of storage into logical units called files.
- Each storage medium is controlled by a device, such as a disk drive or tape drive.
- Different media have varying properties, including access speed, capacity, data-transfer rate, and access method (sequential or random).
- **File-System Management:**
  - Files are usually organized into directories.
  - Access control determines who can access what files.
  - Operating system activities include:
    - Creating and deleting files and directories.
    - Providing functions to manipulate files and directories.
    - Mapping files onto secondary storage.
    - Backing up files onto stable (non-volatile) storage media.

## I/O Subsystem

- One purpose of OS is to hide peculiarities of hardware devices from the user
- I/O subsystem responsible or
  - Memory management of I/O including
    - buffering (storing data temporarily while it is being transferred),
    - caching (storing parts of data in faster storage for performance),
    - spooling (the overlapping of output of one job with input of other jobs)
  - General device-driver interface
  - Drivers for specific hardware devices

## Protection and Security

- **Protection** – any mechanism for controlling access of processes or users to resources defined by the OS
- **Security** – defense of the system against internal and external attacks
  - Huge range, including denial-of-service, worms, viruses, identity theft, theft of service

Systems generally first distinguish among users, to determine who can do what

- User identities (**user IDs**, security IDs) include name and associated number, one per user
- User ID then associated with all files, processes of that user to determine access control
- Group identifier (**group ID**) allows set of users to be defined and controls managed, then also associated with each process, file
- **Privilege escalation** allows user to change to effective ID with more rights

## **Types of operating system**

- Many different **types of operating systems** are involved till date. The operating systems are improved in terms of their capabilities.
- The modern days operating systems allows multiple users to carry out multiple tasks simultaneously.
- Based on their capabilities and the types of application supported, the operating systems can be divided into following six major categories.
  1. **Batch processing operating systems**
  2. **Multi user operating systems**
  3. **Multitasking operating systems**
  4. **Real time operating systems**
  5. **Multi processors operating systems**
  6. **Embedded operating systems**

## **Batch Processing Operating Systems**

- A batch processing operating system executes one job at a time.
- Jobs are combined into batches and submitted to the system as input data.
- Jobs in batches are processed on a first-come, first-served basis.
- After completing one job, the operating system fetches the next job from the input data.
- There is no need for human interaction before fetching the next job after completion.
- Examples of batch processing operating systems include:

**IBM OS/360, VAX/VMS.**

## **Multi user operating systems**

- A multi-user operating system allows multiple users to access and use a single computer system simultaneously.
- These operating systems are specifically designed for multi-user environments, enabling efficient resource sharing.
- Examples of multi-user operating systems include:

**Unix, Linux, Windows 2000, M-386**

## **Multitasking Operating Systems**

- A multitasking operating system allows a user to perform multiple tasks simultaneously on a single computer system.
- It is also known as a multiprocessing operating system or multiprogramming operating system.
- The first multitasking operating system was created in the 1960s.
- The number of tasks that can be processed simultaneously depends on the speed of the CPU, the capacity of memory, and the size of the programs.
- Examples of multitasking operating systems include:

**Linux, Unix, Windows 2000, Windows XP, Windows 10**

## **Multiprocessor Operating Systems**

- A multiprocessor operating system allows a computer system to use more than one CPU to execute multiple processes simultaneously.
- Systems with multiple CPUs can process tasks faster than those with a single CPU.
- This type of operating system is designed to manage and coordinate the use of multiple processors effectively.
- Examples of multiprocessor operating systems include:

**Linux, Unix, Windows 2000**

## **Embedded Operating Systems**

- An embedded operating system is similar to a real-time operating system, designed specifically for embedded computer systems.
- These operating systems are primarily used to perform computational tasks in electronic devices.
- Embedded systems are often found in devices such as appliances, automotive controls, and medical equipment.
- Examples of embedded operating systems include:

**Palm Operating System, Windows CE**

## Real-Time Operating Systems

- A real-time operating system is similar to a multitasking operating system but is specifically designed to handle real-time applications.
- Real-time applications must execute within a specific time frame, making time a critical constraint.
- Examples of real-time applications include robots and machine learning systems.
- There are two main types of real-time operating systems:
  - **Hard Real-Time Operating Systems:** These systems require strict adherence to timing constraints; missing a deadline can result in failure.
  - **Soft Real-Time Operating Systems:** These systems are more flexible with timing; occasional delays are acceptable but should be minimized.
- Examples of real-time operating systems include: **MTOS, Lynx, RTX.**

### Hard Real-Time Systems:

- These systems are purely deterministic, meaning they must meet strict deadlines.
- For example, if a user expects output within 10 seconds, the system must deliver it exactly at the 10-second mark.
- Missing the deadline (e.g., delivering at 9 or 11 seconds) results in system failure.
- An example is a missile defense system, where precise timing is critical; if a missile is supposed to hit a target at a specific time, any delay could lead to failure.

### Soft Real-Time Systems:

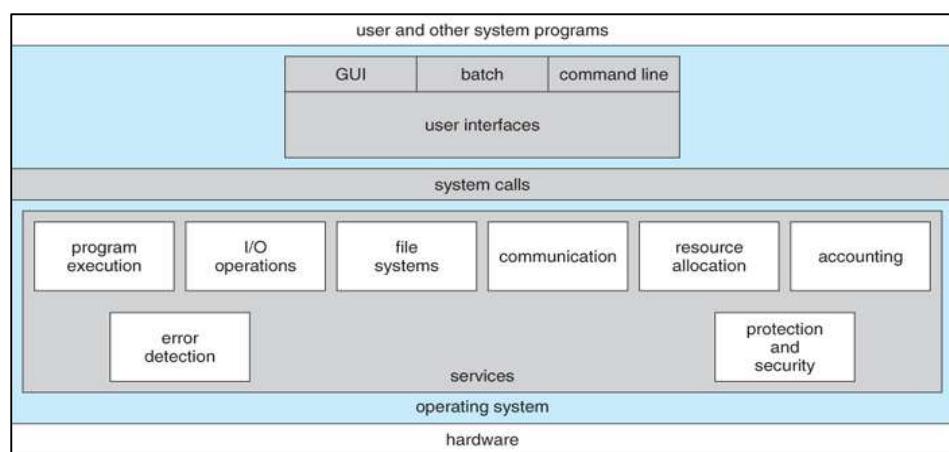
- In these systems, meeting deadlines is not always mandatory for every task.
- While tasks should ideally complete on time, occasional deadline misses are acceptable.
- However, if deadlines are consistently missed, the system's performance may degrade significantly.
- Examples include personal computers and audio/video systems, where some delays can be tolerated without critical consequences.

## Distributed Systems

- A distributed system is a collection of separate, possibly heterogeneous, systems that are networked together.
- The network serves as a communication path, with TCP/IP being the most common protocol used.
- Types of networks include:
  - Local Area Network (LAN)
  - Wide Area Network (WAN)
  - Metropolitan Area Network (MAN)
  - Personal Area Network (PAN)
- A network operating system provides features that facilitate interaction between systems across the network.
- Communication schemes allow these systems to exchange messages seamlessly.
- Users experience the illusion of a single coherent system despite the underlying complexity.

## OS Structure: Layered Approach

- The operating system is structured in multiple layers, each built on top of lower layers.
- The bottom layer (Layer 0) consists of the hardware, while the top layer (Layer N) is the user interface.
- Each layer utilizes functions and services from only the lower-level layers, promoting modularity.
- This layered approach simplifies system design and enhances maintainability.



## 1. Monolithic Structure:

In a monolithic OS, all core functions (e.g., file systems, device drivers, system service calls) operate within the kernel space. Every component communicates directly, which boosts efficiency by eliminating communication overhead.

- **Advantages:** Fast due to direct communication between components.
- **Disadvantages:** A bug in one module can crash the entire OS; complex to modify due to tight integration.
- **Examples:** Unix, Linux (although modern Linux has modular aspects, it remains primarily monolithic), MS-DOS.

## 2. Microkernel Structure:

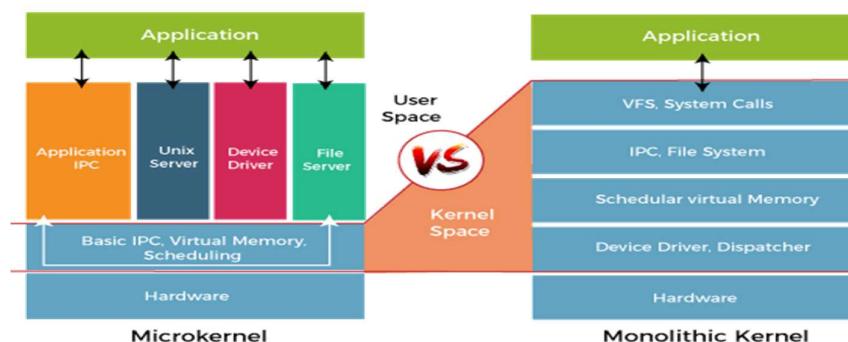
Only the basic system components, such as inter-process communication, scheduling, and memory management, run in the kernel space. All other services (device drivers, file systems) run in user space and interact with the kernel via message passing.

- **Advantages:** Improved modularity and reliability, as failure in one service won't crash the whole system; easier to maintain and extend.
- **Disadvantages:** Potential performance overhead due to increased communication between user-space and kernel-space services.
- **Examples:** QNX, Minix.

## 3. Hybrid Structure:

Combines elements of both monolithic and microkernel designs. Essential services run in the kernel (like in monolithic systems), while other services run in user space (like in microkernel systems). It seeks to balance performance and modularity.

- **Advantages:** Performance is closer to monolithic systems, with improved modularity and reliability.
- **Disadvantages:** Can still suffer from performance trade-offs, but generally more efficient than pure microkernels.
- **Examples:** Windows NT (and successors like Windows 2000, XP, 7, 10), macOS, BeOS.



## System Calls

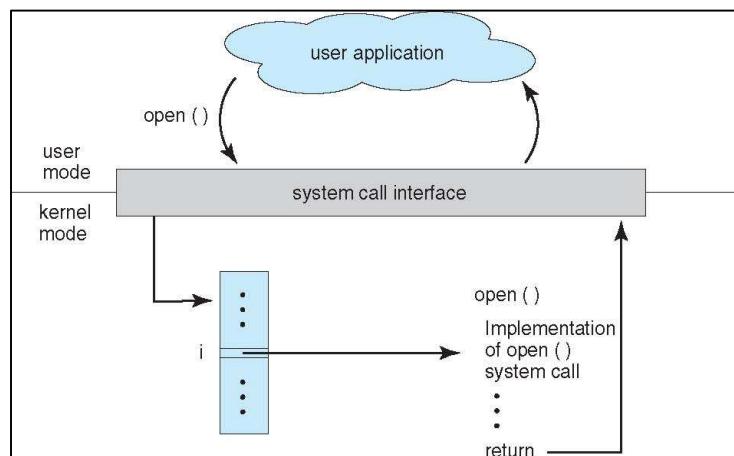
A system call is a mechanism that provides an interface between a process and the operating system, allowing programs to request services from the OS kernel.

- **Programmatic Access:** They enable programs to execute functions such as file manipulation, process control, and communication through a standardized interface.
- **API Integration:** System calls are accessed via high-level Application Programming Interfaces (APIs), simplifying the process for developers to utilize OS services without needing to understand the underlying complexities.
- **Controlled Access:** By using system calls, the OS can enforce security and stability, ensuring that applications interact with hardware and system resources in a safe manner.
- **Examples of Common APIs:**
  - **Win32 API:** Used for Windows applications, providing access to various OS functionalities.
  - **POSIX API:** Standardized interface for UNIX-like systems, facilitating compatibility across different platforms.
  - **Java API:** Allows Java applications to interact with the underlying operating system through the Java Virtual Machine (JVM).

These features make system calls essential for effective communication between user applications and the operating system, ensuring efficient resource management and security.

- **Example:** A program uses the `open()` system call to open a file, then uses `read()` to read data, and finally `close()` to close the file.

```
int fd = open("file.txt", O_RDONLY); // Open the file in read-only mode
read(fd, buffer, sizeof(buffer));   // Read data from the file
close(fd);                         // Close the file
```



## **Types of System Calls**

### **1. Process Control**

- end, abort: Terminate a process.
- load, execute: Load a program into memory and execute it.
- create process, terminate process: Manage process lifecycle.
- get/set process attributes: Retrieve or modify process information.
- wait for time: Pause execution for a specified duration.
- wait event, signal event: Synchronize processes.
- allocate and free memory: Manage memory resources.

### **2. File Management**

- create file, delete file: Manage file creation and deletion.
- open, close file: Access and release files.
- read, write, reposition: Manipulate file data.
- get/set file attributes: Retrieve or modify file information.

### **3. Device Management**

- request device, release device: Manage device access.
- read, write, reposition: Interact with device data.
- get/set device attributes: Retrieve or modify device information.
- logically attach/detach devices: Manage device connections.

### **4. Information Maintenance**

- get/set time or date: Manage system time settings.
- get/set system data: Access or modify system-wide information.
- get/set process, file, or device attributes: Manage various resource attributes.

### **5. Communications**

- create/delete communication connection: Establish or terminate connections between processes.
- send/receive messages: Facilitate inter-process communication.
- transfer status information: Share status updates between processes.
- attach/detach remote devices: Manage connections to remote hardware.

## Examples of Windows and Unix System Calls

	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()

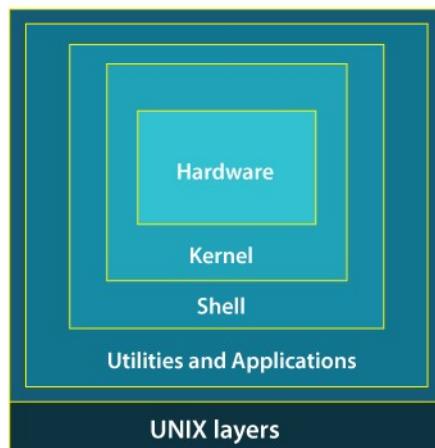
## Unix Architecture

**Layer-1: Hardware:** It consists of all hardware related information.

**Layer-2: Kernel:** This is the core of the Operating System. It is a software that acts as the interface between the hardware and the software. Most of the tasks like memory management, file management, network management, process management, etc., are done by the kernel.

**Layer-3: Shell commands:** This is the interface between the user and the kernel. Shell is the utility that processes your requests. When you type in a command at the terminal, the shell interprets the command and calls the program that you want. There are various commands like cp, mv, cat, grep, id, wc, nroff, a.out and more.

**Layer-4: Application Layer:** It is the outermost layer that executes the given external applications.



## **System Kernel of Unix**

The **System Kernel** is the core component of the Unix operating system. It acts as the bridge between the hardware and user applications, managing system resources, processes, and communication. Below is a brief overview of the main components of the Unix kernel, accompanied by a simplified sketch to illustrate its structure. Main Tasks of the Kernel

### **1. Process Management:**

- Handles the creation, scheduling, and termination of processes.
- Manages CPU allocation and process synchronization.

### **2. Device Management:**

- Controls and coordinates access to hardware devices.
- Manages device drivers for input/output operations.

### **3. Memory Management:**

- Allocates and deallocates memory for processes.
- Implements virtual memory to extend physical memory usage.

### **4. Interrupt Handling:**

- Responds to hardware interrupts to manage asynchronous events.
- Ensures timely processing of critical tasks.

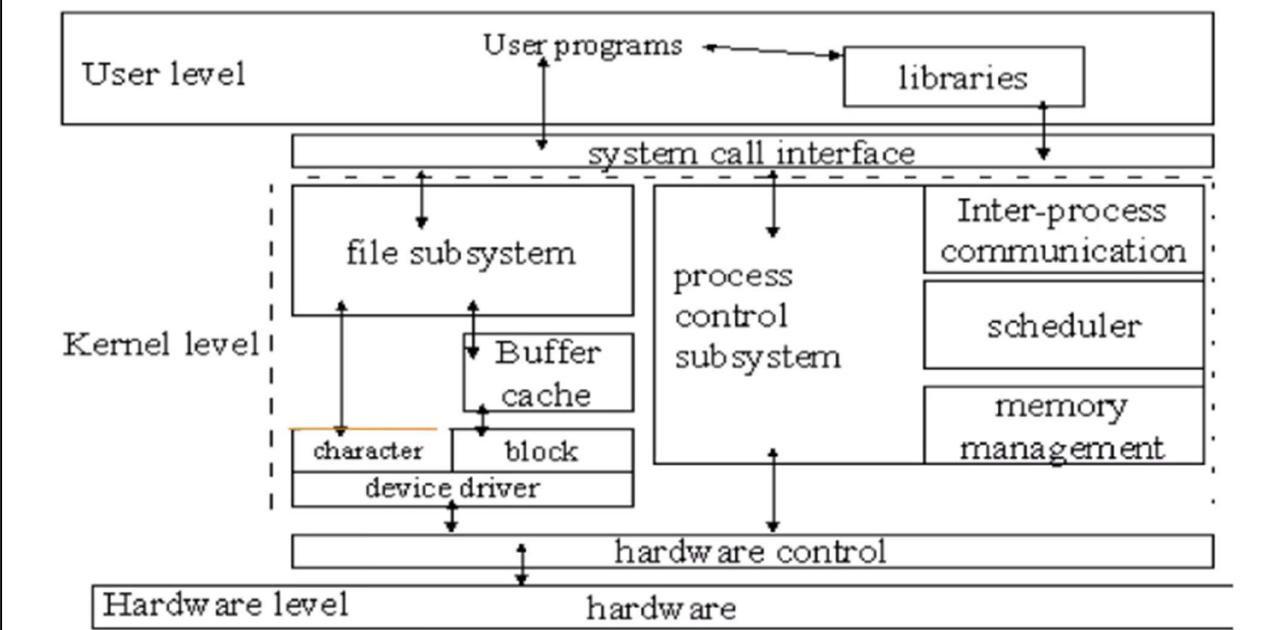
### **5. I/O Communication:**

- Manages data transfer between devices and applications.
- Provides a standardized interface for I/O operations.

### **6. File System Management:**

- Organizes and manages files on storage devices.
- Provides services for file creation, deletion, reading, and writing.

## Structure: UNIX Kernel



### Explanation of the Sketch:

- **User Space:** This is where user applications run. Users interact with the system through applications (like the shell) that invoke system calls to request services from the kernel.
- **System Calls:** User applications interact with the kernel through system calls, allowing them to perform operations like creating processes, accessing files, and managing memory.
- **Kernel Space:** The kernel operates in a protected space with direct access to hardware. This is where process management, memory management, file system management, device drivers, and IPC are handled.
- **Hardware:** At the bottom, the kernel interfaces with the physical hardware of the computer, including CPUs, memory, and I/O devices.

**Q. Define operating system. Explain RTOS services in contrast with traditional OS.**

Feature	RTOS (Real-Time Operating System)	Traditional OS
Purpose	Designed for real-time applications with strict deadlines.	General-purpose computing for various user tasks.
Task Scheduling	Deterministic scheduling to meet real-time constraints.	Fair scheduling, prioritizing tasks without strict time limits.
Response Time	Predictable and low-latency response times.	Higher latency; response time is less critical.
Multitasking	Preemptive multitasking, prioritizing time-sensitive tasks.	Multitasking based on priority and availability but not real-time.
Reliability	High reliability, suitable for mission-critical tasks.	Reliable, but not designed for continuous or failure-intolerant systems.
Examples	VxWorks, FreeRTOS, QNX	Windows, Linux, macOS
Use Cases	Embedded systems (medical devices, avionics).	Desktop computers, servers, mobile devices.

**Q. One of the important functionalities of an operating system is process management. State the system calls that are used to manage the process with the input parameters (5 system calls are expected, 2 marks for each)**

System Call	Description	Input Parameters
<code>fork()</code>	Creates a new process by duplicating the calling process.	None (inherits memory and context from the parent).
<code>exec()</code>	Replaces the current process image with a new program.	<code>const char *filename</code> , <code>char *const argv[]</code> , <code>char *const envp[]</code> .
<code>wait()</code>	Waits for a child process to terminate.	<code>int *status</code> (pointer to store the child's exit status).
<code>exit()</code>	Terminates the calling process.	<code>int status</code> (exit status code for the parent process).
<code>getpid()</code>	Returns the process ID of the calling process.	None (returns the PID as an integer).

# Process Management

## Process

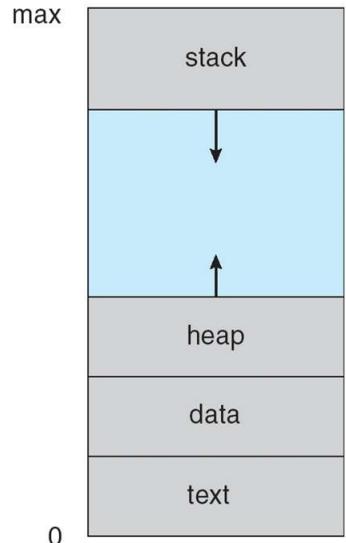
A **process** is a program in execution. It represents the basic unit of work in a computer system. The execution of a process progresses in a **sequential fashion**, meaning that instructions are executed one after another.

- A **program** is a passive entity (stored on a disk as an executable file), whereas a **process** is an active entity that is loaded into memory for execution.
- A program becomes a process when its executable file is loaded into memory and starts execution.
- The execution of a program can be initiated through user actions like GUI mouse clicks or command line inputs.

## Process Components:

### 1. Stack:

- Contains temporary data such as function/method parameters, return addresses, and local variables used during execution.



### 2. Heap:

- Dynamically allocated memory used by the process during its runtime.

### 3. Text (Code Segment):

- Contains the executable instructions and the current activity represented by the value of the program counter and the contents of processor registers.

### 4. Data Section:

- Stores global and static variables that the process can access throughout its execution.

## Process States

### 1. Start:

- This is the initial state when the process is first created or started.
- The operating system initializes the process.

### 2. Ready:

- The process is waiting to be assigned to a CPU for execution.
- It may enter this state after being started or while running, if interrupted by the scheduler to allocate CPU time to another process.

### 3. Running:

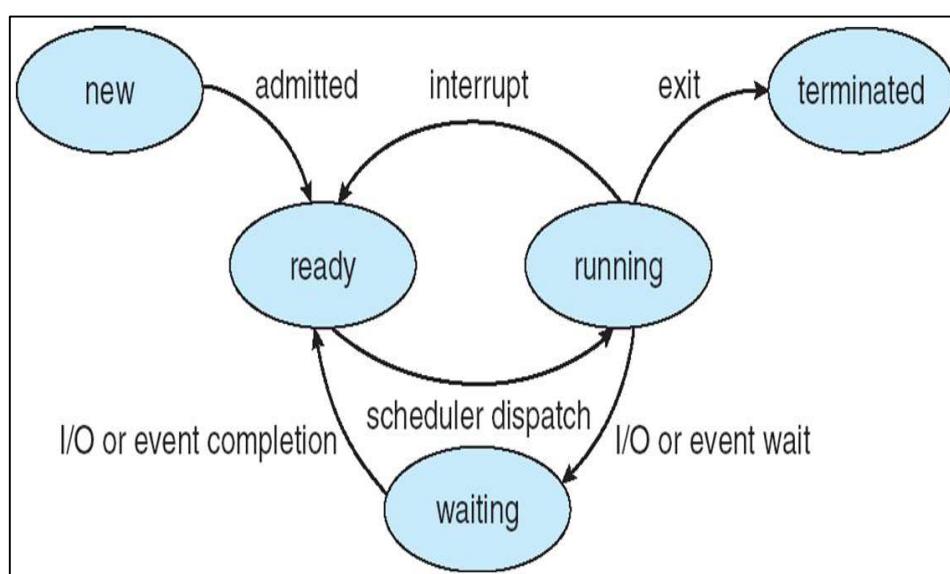
- When the process is assigned to a CPU by the OS scheduler, it transitions to the running state.
- The process instructions are actively executed by the processor.

### 4. Waiting:

- The process enters the waiting state if it requires a resource (such as waiting for I/O operations or user input) and cannot proceed until the resource is available.

### 5. Terminated/Exit:

- Once the process completes its execution or is forcibly terminated by the operating system, it moves to the terminated state.
- In this state, the process is removed from memory.



## Process Control Block (PCB)

The **Process Control Block (PCB)** is a data structure maintained by the operating system for each process. It stores essential information required to manage the process during its lifecycle.

### Information Contained in a PCB:

#### 1. Process State:

- The current state of the process (e.g., running, waiting, ready).

#### 2. Program Counter:

- Stores the address of the next instruction to be executed for the process.

#### 3. CPU Registers:

- Contains the values of the CPU registers specific to the process, which need to be saved when the process is not running and restored when the process is resumed.

#### 4. CPU Scheduling Information:

- Includes process priorities and pointers to scheduling queues, helping the operating system decide which process to execute next.

#### 5. Memory Management Information:

- Information about the memory allocated to the process, such as base and limit registers, page tables, or segment tables.

#### 6. Accounting Information:

- Tracks the amount of CPU time the process has used, the elapsed clock time since the process started, and any time limits imposed.

#### 7. I/O Status Information:

- Lists the I/O devices allocated to the process and the status of open files.

process state
process number
program counter
registers
memory limits
list of open files
• • •

## Context Switch

A **context switch** occurs when the CPU switches from one process or thread to another. This mechanism is essential for multitasking in operating systems, enabling the CPU to handle multiple processes without requiring additional processors.

### When is Context Switching Required?

#### 1. Multitasking:

- When the CPU needs to switch between different processes, moving them in and out of memory.

#### 2. Kernel/User Mode Switch:

- Sometimes occurs when switching between user applications and kernel (OS) services.

#### 3. Interrupts:

- Triggered when the CPU is interrupted by external events, such as I/O operations, which require the CPU to handle those events.

### Steps in a Full Context Switch:

1. **Save the Context:** The CPU saves the state of the running process, including the program counter and register values.
2. **Update the PCB:** The Process Control Block (PCB) of the current process is updated with the reason for stopping and other important state information, such as memory, CPU usage, and process state.
3. **Move Process to Queue:** The current process is moved to the appropriate queue (e.g., ready, blocked) depending on its state.
4. **Select New Process:** The CPU scheduler selects a new process from the ready queue for execution.
5. **Update the New Process's PCB:** The PCB of the selected process is updated to reflect that it is now in the **Running** state.
6. **Update Memory Management:** If necessary, update memory management structures to map the selected process's address space.
7. **Restore Context:** The CPU restores the saved state (registers, program counter) of the new process so it can continue executing from where it last stopped.

## Process Scheduling

Process scheduling is a core function of the operating system, designed to maximize the use of the CPU by switching between processes efficiently. This allows for time-sharing, where multiple processes can appear to be running simultaneously.

### Key Points:

- **Maximize CPU Utilization:**  
The goal of process scheduling is to keep the CPU busy at all times by selecting processes to execute in rapid succession.
- **Process Scheduler:**  
The scheduler is responsible for choosing the next process to be executed from a set of processes waiting to run.
- **Scheduling Queues:**
  1. **Job Queue:**  
This contains all the processes in the system, including those that are not yet in the ready state.
  2. **Ready Queue:**  
A list of processes that are loaded in main memory and are ready to be executed by the CPU.
  3. **Device Queues:**  
Processes waiting for specific I/O operations to complete (e.g., waiting for a printer or disk).

## Schedulers

### 1. Short-term Scheduler (CPU Scheduler):

- Decides which process should be executed next and allocates the CPU to it.
- It is invoked frequently (every few milliseconds), so it must be fast.
- Sometimes the only scheduler in a system.

### 2. Long-term Scheduler (Job Scheduler):

- Determines which processes should be loaded into the **ready queue** from the job pool.
- It is invoked less frequently (seconds or minutes), so it can be slower.
- Controls the degree of **multiprogramming** (the number of processes in memory).
- Strives to balance between **I/O-bound** processes (perform many short CPU bursts) and **CPU-bound** processes (perform longer computations).

### 3. Medium-term Scheduler:

- This scheduler is used if the system needs to reduce the degree of multiprogramming.
- It temporarily removes processes from memory and places them on disk (swapping).
- The process can be brought back into memory to resume execution when needed.

Aspect	Long-Term Scheduler	Short-Term Scheduler	Medium-Term Scheduler
Function	Job scheduler	CPU scheduler	Process swapping scheduler
Speed	Slower than the short-term scheduler	Fastest among all schedulers	Speed is between short and long-term
Control over Multiprogramming	Controls the degree of multiprogramming	Provides lesser control over multiprogramming	Reduces the degree of multiprogramming
Presence in Time Sharing Systems	Minimal or almost absent in time-sharing systems	Minimal in time-sharing systems	Present in time-sharing systems
Selection Criteria	Selects processes from the job pool and loads them into memory	Selects processes that are ready for execution	Reintroduces processes into memory for continued execution

## Operations on Processes

### 1. Process Creation

- **Hierarchy:** Processes are created in a parent-child relationship, forming a tree structure.
- **Process Identifier (PID):** Each process is identified and managed using a unique process identifier (PID).
- **Resource Sharing Options:**
  - **All Resources:** The parent and children share all resources.
  - **Subset of Resources:** The children share only a subset of the parent's resources.
  - **No Resources:** The parent and children share no resources at all.
- **Execution Options:**
  - **Concurrent Execution:** Parent and children can execute concurrently.
  - **Synchronous Execution:** The parent waits for children to terminate before continuing.

### 2. Process Termination

- **Termination Process:**
  - A process executes its last statement and requests deletion through the exit() system call.
  - Returns status data from the child to the parent via the wait() system call.
  - The operating system deallocates the process's resources.
- **Parent Termination of Child Processes:**
  - The parent can forcibly terminate child processes using the abort() system call for various reasons, such as:
    - The child has exceeded allocated resources.
    - The task assigned to the child is no longer necessary.
    - The parent is terminating, and the operating system does not allow the child to continue without its parent.

## Models of Inter-Process Communication (IPC)

### 1. Shared Memory

In the shared memory model, multiple processes are allowed to access the same region of memory. This memory segment is set up by one process, and other processes can map this memory into their address space.

#### Characteristics:

- **Speed:** Shared memory is typically faster than other IPC methods because it allows processes to read and write directly to the memory area without involving the kernel for every communication. The overhead of system calls is minimized.
- **Synchronization:** Since multiple processes can read from and write to shared memory simultaneously, synchronization mechanisms (like semaphores, mutexes, or monitors) are necessary to prevent race conditions and ensure data consistency. Without synchronization, data corruption or unpredictable results may occur.

**Use Cases:** Applications that require high-speed communication and share large amounts of data, such as multimedia processing, game engines, or scientific simulations.

### 2. Message Passing

In the message passing model, processes communicate by sending and receiving messages. This can occur in two ways: synchronous or asynchronous communication.

#### Characteristics:

- **Structure:** The message passing model is more structured compared to shared memory. It provides a clear interface for communication, where data is packaged into messages and sent between processes.
- **Ease of Implementation:** This model can be easier to implement, especially in distributed systems, as it abstracts the underlying details of how data is shared. Each process can operate independently, without needing to worry about managing shared memory.
- **Performance:** Message passing can be slower than shared memory due to the overhead of copying messages and potential context switches. However, it can be more flexible and robust, especially in systems where processes do not reside on the same machine.

**Use Cases:** Distributed systems, microservices architectures, and client-server applications, where processes may run on different machines or where isolation between processes is desirable for security or stability.

## Synchronization in Inter-Process Communication (IPC)

**Synchronization** is vital for coordinating processes that communicate through `send()` and `receive()` calls. Here's a brief overview of the different types:

### Message Passing Synchronization

#### 1. Blocking Send:

- The sender waits until the message is received.
- **Use:** Ensures reliable communication.

#### 2. Non-Blocking Send:

- The sender sends the message and continues without waiting.
- **Use:** Increases throughput and concurrency.

#### 3. Blocking Receive:

- The receiver waits until a message is available.
- **Use:** Ensures timely processing of required data.

#### 4. Non-Blocking Receive:

- The receiver checks for a message and continues if none is available.
- **Use:** Keeps the receiver responsive.

### Combinations of Send and Receive

- **Blocking Send + Blocking Receive:** Both processes synchronize at the message exchange (rendezvous).
- **Blocking Send + Non-Blocking Receive:** The sender waits while the receiver may continue other tasks.
- **Non-Blocking Send + Blocking Receive:** The sender proceeds while the receiver waits for a message.
- **Non-Blocking Send + Non-Blocking Receive:** Both processes operate independently.

### Producer-Consumer Problem

In the producer-consumer problem, using **blocking send** and **blocking receive** simplifies synchronization. The producer waits for the consumer to receive data before producing more, preventing resource overflow.

## Buffering in Inter-Process Communication

Buffering is essential in inter-process communication (IPC) as it provides a temporary storage mechanism for messages exchanged between processes. The implementation of message queues can be categorized into three types based on their capacity:

### 1. Zero Capacity:

- **Description:** The queue has no capacity for messages.
- **Behavior:** The sender must block (wait) until the recipient receives the message.
- **Use Case:** Suitable for immediate communication where no delays are acceptable.

### 2. Bounded Capacity:

- **Description:** The queue has a finite length  $n$ , allowing at most  $n$  messages to reside in it.
- **Behavior:**
  - If the queue is not full, messages can be sent and the sender can continue executing without waiting.
  - If the queue is full, the sender blocks until space is available.
- **Use Case:** Balances resource use and communication efficiency, ideal for scenarios where message flow can be anticipated.

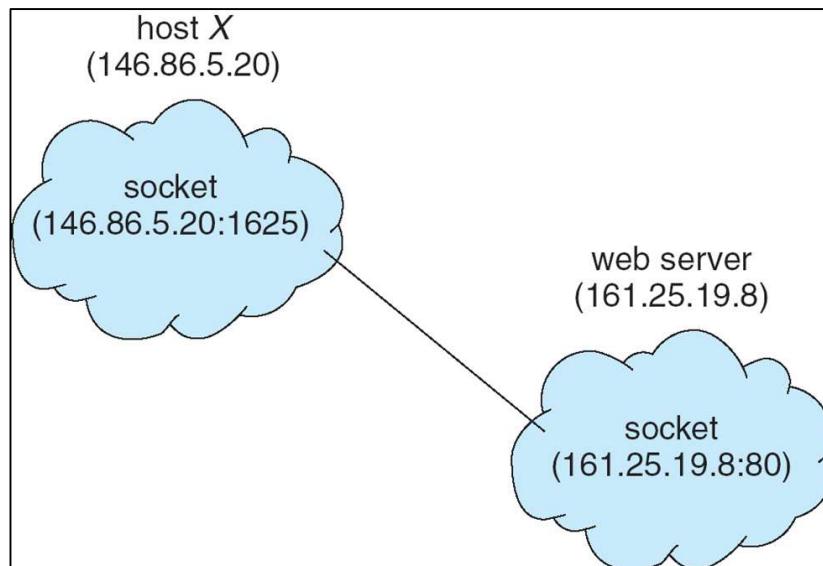
### 3. Unbounded Capacity:

- **Description:** The queue has potentially infinite capacity.
- **Behavior:** The sender never blocks, allowing any number of messages to wait in the queue.
- **Use Case:** Useful in scenarios where message traffic is unpredictable and the system can handle variable loads without performance penalties.

## Sockets in Networking

A socket is an endpoint for communication between two devices on a network.

- **Structure:** A socket is formed by the concatenation of an **IP address** and a **port number**. This combination allows for the differentiation of network services on a host.
  - **Example:** The socket **161.25.19.8:1625** specifies **port 1625** on the host **161.25.19.8**.
- **Communication:** Communication in networking typically occurs between pairs of sockets. Each socket pair represents a unique connection.
- **Port Numbers:**
  - **Well-Known Ports:** Ports below **1024** are designated as well-known ports, reserved for standard services (e.g., HTTP on port 80, HTTPS on port 443).
- **Loopback Address:** The special IP address **127.0.0.1** is known as the loopback address and refers to the local machine. This allows processes to communicate with themselves for testing and development purposes.

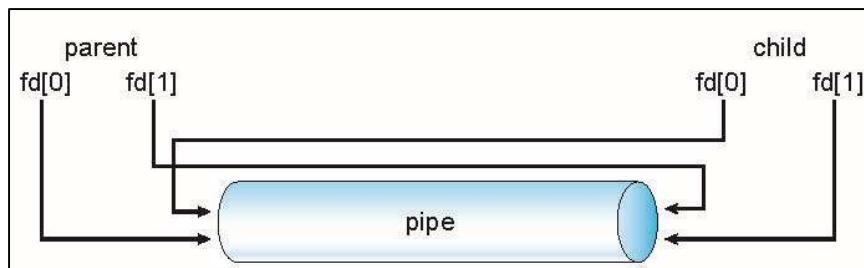


## Pipes in Inter-Process Communication

**Pipes** are a fundamental mechanism in operating systems for facilitating inter-process communication (IPC). They allow one process to send data to another, effectively acting as a conduit. Pipes enable processes to communicate in a producer-consumer fashion, where one process produces data and another consumes it.

### 1. Ordinary Pipes

- **Function:** Facilitate communication in a standard producer-consumer model.
  - **Producer** writes to one end (write-end).
  - **Consumer** reads from the other end (read-end).
- **Characteristics:**
  - **Unidirectional:** Data flows in one direction only (from producer to consumer).
  - **Parent-Child Relationship:** Typically require a relationship between the communicating processes, where one process (usually the parent) creates the pipe and communicates with its child process.
- **Platform:**
  - Commonly referred to as **anonymous pipes** in Windows systems.



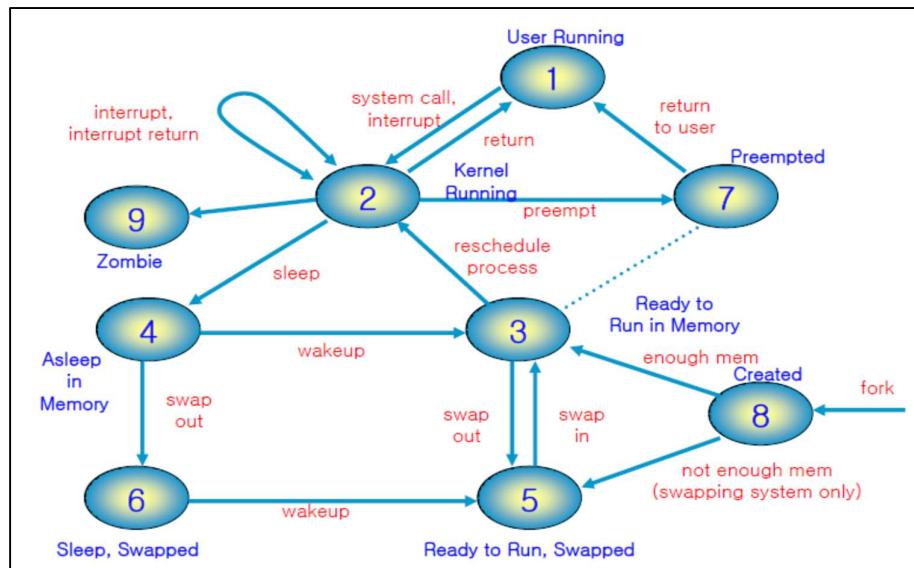
### 2. Named Pipes

- **Function:** Provide a more versatile communication method than ordinary pipes.
- **Characteristics:**
  - **Bidirectional:** Allows two-way communication between processes.
  - **No Parent-Child Relationship:** Can be accessed by processes that do not have a parent-child relationship.
  - **Multiple Processes:** Several processes can use the same named pipe for communication.
- **Platform:** Available on both UNIX and Windows systems.

## Unix-Process States

The lifecycle of a Unix process can be represented through various states, each describing the process's current status. Below is a brief explanation of these states along with their transitions:

1. **User Mode:** The process is actively executing user-level code, performing tasks assigned to it.
2. **Kernel Mode:** The process is executing system calls or kernel-level code, allowing it to interact with the hardware and access protected resources.
3. **Ready:** The process is not currently executing but is ready to run. It is waiting for the kernel to schedule it on the CPU.
4. **Sleeping (In Memory):** The process is in a sleeping state, often waiting for an event or resource to become available while still residing in main memory.
5. **Swapped Ready:** The process is ready to run, but it is currently swapped out of main memory. The swapper must load it back into memory before it can execute.
6. **Swapped Sleeping:** The process is in a sleeping state but has been swapped out to secondary storage to free up memory for other processes.
7. **Preempted:** The process is returning from kernel mode to user mode, but the kernel interrupts it to perform a context switch, scheduling another process for execution.
8. **Newly Created:** The process has just been created and exists in a transition state. It is not ready to run yet and is neither in a sleeping state nor executing.
9. **Zombie:** The process has finished executing and called the exit system call. Although it no longer exists as a running process, it leaves behind a record (exit code and statistics) for its parent process to collect. This is the final state of a process.



## Program vs. Process

- **Program:**

A program is a set of instructions written by a programmer using a programming language. It is a static, passive entity stored on disk that contains the logic to solve a specific task. A program does not perform any action until it is executed.

- **Process:**

A process is an active, dynamic instance of a program that is being executed. It is created when a program is loaded into memory and starts running. A process includes the program code as well as its current activity (program counter, registers, etc.).

### Key Difference:

- A **program** is static (passive) code, while a **process** is a dynamic (active) execution of that code.

## Multiprocess Architecture in Google Chrome

### 1. Browser Process:

- Manages the user interface and handles disk/network I/O.
- Responsible for user interactions and tab management.

### 2. Renderer Process:

- Renders web pages and processes HTML, CSS, and JavaScript.
- A new process is created for each website, isolating tabs for stability.
- Runs in a sandbox to restrict access and enhance security.

### 3. Plug-in Process:

- Dedicated process for each type of plug-in.
- Operates independently to prevent crashes affecting the browser.

# Multithread Programming and CPU Scheduling

## What is a Thread?

- A thread is a **flow of execution** through a process with its own program counter, system registers, and stack.
- Threads share certain resources with other threads in the same process, such as the **code segment, data segment, and open files**.

## Characteristics of Threads

- Threads are also known as **lightweight processes**.
- They enable **parallelism**, improving application performance.
- Provide a **software-based approach** to enhance OS performance by reducing overhead.
- A thread is equivalent to a classical process but exists within a parent process.
- Each thread represents an independent flow of control.
- Frequently used in **network servers** and **web servers**.
- Suitable for parallel execution on **shared memory multiprocessors**.

## Advantages of Threads

- **Minimized Context Switching:** Less overhead compared to processes.
- **Concurrency:** Threads run parallelly within the same process.
- **Efficient Communication:** Threads share memory space, facilitating quick data sharing.
- **Economical:** Creating and switching threads is resource-efficient.
- **Scalability:** Threads maximize the efficiency of multiprocessor systems.

## Key Benefits

- **Responsiveness:** Enhances application responsiveness by enabling concurrent execution.
- **Resource Sharing:** Threads share the same process resources efficiently.
- **Economy:** Reduces the cost of creating and managing execution units.
- **Scalability:** Provides a robust foundation for parallelism in multicore systems.

## Single-threaded & Multithreaded processes

### 1. Single-Threaded Process:

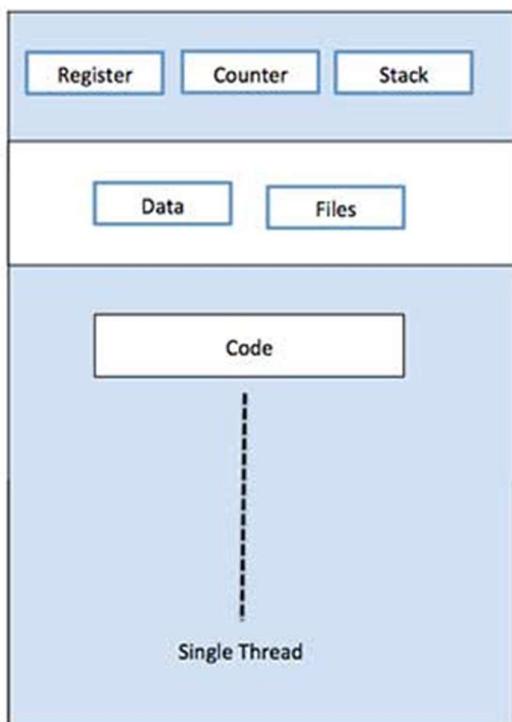
- A single-threaded process has only one flow of execution.
- Tasks are executed sequentially, one after another.
- If one task blocks (e.g., a system call), the entire process is halted until the blocking task is completed.

**Example:** A basic text editor that processes input, saves files, or formats text one task at a time.

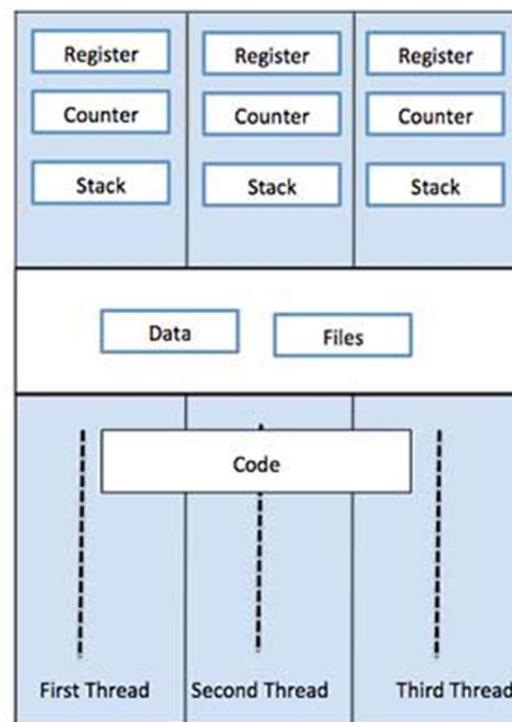
### 2. Multithreaded Process:

- A multithreaded process allows multiple threads to execute concurrently.
- Each thread represents a separate task or operation.
- Threads within a process share memory and resources but execute independently.
- Even if one thread is blocked, other threads can continue execution.

**Example:** A web browser with threads for rendering pages, handling user input, and managing downloads simultaneously.



Single Process P with single thread

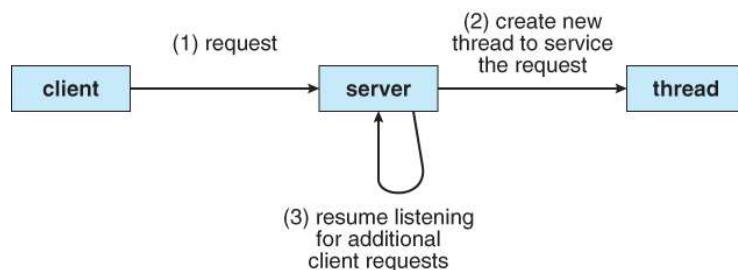


Single Process P with three threads

## Comparison Between Processes and Threads

Aspect	Process	Thread
Weight	Heavyweight and resource-intensive.	Lightweight and uses fewer resources.
Switching	Requires interaction with the operating system, adding overhead.	Does not require OS interaction, making switching faster.
Memory and Resources	Each process has its own memory and file resources.	Threads share memory and resources within the same process.
Blocking	If one process is blocked, no other process can execute until it is unblocked.	A blocked thread does not halt other threads in the same process.
Resource Utilization	Multiple processes consume more resources compared to threads.	Multiple threads consume fewer resources.
Independence	Processes operate independently of one another.	Threads can interact and modify each other's data.
Use of Open Files	Open files are not shared between processes.	All threads share the same set of open files and child processes.
Execution in Parallel Environments	Each process executes its own code independently.	Threads can run concurrently, allowing tasks to overlap effectively.

## Multithreaded Server Architecture



- Benefits

- **Responsiveness:** Allows continued execution even if part of the process is blocked, which is especially important for user interfaces.
- **Resource Sharing:** Threads share the resources of the process, making it easier to manage than using shared memory or message passing.
- **Economy:** Creating threads is cheaper than creating processes, and thread switching has lower overhead compared to context switching.
- **Scalability:** A multithreaded server can take advantage of multiprocessor architectures, improving performance and handling more tasks concurrently.

## **Types of Threads**

### **1. User Level Threads**

These are managed entirely in user space without kernel support. The thread library handles the creation, destruction, communication, scheduling, and context management of threads.

- **Advantages:**

- Thread switching does not require kernel mode privileges.
- User level threads can run on any operating system.
- Scheduling can be application-specific.
- User level threads are faster to create and manage.

- **Disadvantages:**

- Most system calls are blocking, which can affect thread execution.
- Multithreaded applications cannot take full advantage of multiprocessing.

### **2. Kernel Level Threads**

Managed and supported directly by the operating system. The kernel handles thread creation, scheduling, and management.

- **Advantages:**

- The kernel can schedule multiple threads from the same process on different processors.
- If one thread in a process is blocked, the kernel can schedule another thread from the same process.
- Kernel routines themselves can be multithreaded.

- **Disadvantages:**

- Kernel threads are generally slower to create and manage than user threads.
- A mode switch to the kernel is required when transferring control between threads in the same process.

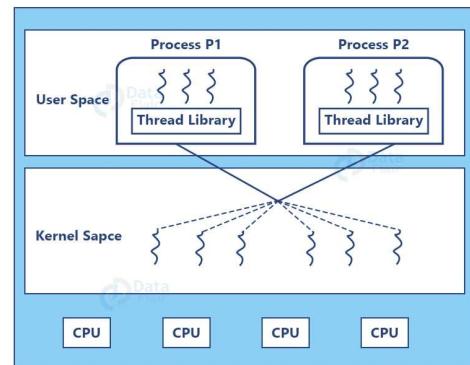
## Difference between User-Level and Kernel-Level Threads

S.N .	User-Level Threads	Kernel-Level Thread
1	User-level threads are faster to create and manage.	Kernel-level threads are slower to create and manage.
2	Implementation is by a thread library at the user level.	Operating system supports creation of Kernel threads.
3	User-level thread is generic and can run on any operating system.	Kernel-level thread is specific to the operating system.
4	Multi-threaded applications cannot take advantage of multiprocessing.	Kernel routines themselves can be multithreaded

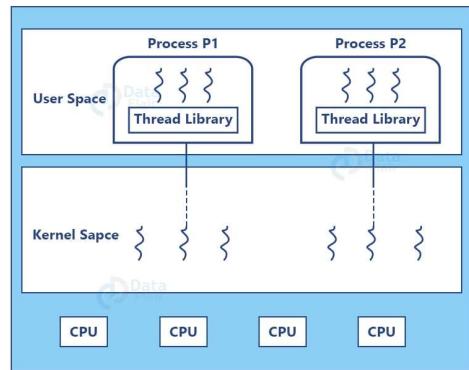
## Multithreading Models

Some operating systems provide a combined user-level thread and kernel-level thread facility, with Solaris being a prominent example. This model allows multiple threads within the same application to run in parallel on multiple processors, and a blocking system call does not block the entire process.

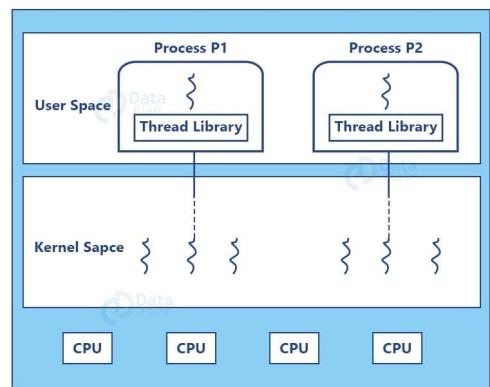
- Many-to-Many Model:** In this model, multiple user-level threads are multiplexed onto an equal or smaller number of kernel-level threads. This allows the system to manage many user threads while utilizing multiple processor cores for better concurrency. It enables parallel execution and the scheduling of threads when a blocking system call occurs.



2. **Many-to-One Model:** Many user-level threads map to a single kernel-level thread. Thread management is handled by a thread library in user space. If a thread makes a blocking system call, the entire process is blocked. Only one thread can access the kernel at a time, and this model does not support parallel execution on multiprocessors.



3. **One-to-One Model:** In this model, there is one kernel-level thread for each user-level thread. This model offers better concurrency than the many-to-one model and supports parallel execution on multiprocessors. However, creating a user thread requires the creation of a corresponding kernel thread.



## Thread Libraries

**Thread Libraries** provide APIs for creating and managing threads. They can be implemented either in user space or kernel space. Some common thread libraries include:

- **POSIX Pthreads:** Available on systems like Linux and macOS, providing a standard thread API. The implementation varies by system, but the API remains consistent across platforms.
- **Win32 Threads:** Kernel-managed threads on Windows with syntax differences compared to POSIX threads. They are created and managed by the Windows kernel.
- **Java Threads:** Java uses threads by default, with implementation depending on the OS. It uses POSIX Pthreads on UNIX systems and Win32 threads on Windows, managed by the Java Virtual Machine (JVM).

## Threading Issues

- **Semantics of fork() and exec():**
  - **fork()**: When called, it creates a new process. In multithreaded programs, it's important to decide if all threads or just the calling thread should be duplicated. Systems offer different versions of fork(), like fork1() (duplicates the calling thread) and forkall() (duplicates all threads).
  - **exec()**: Replaces the current process image with a new one, terminating all threads in the process.
- **Thread Cancellation:**
  - **Asynchronous Cancellation**: Terminates a thread immediately, potentially in the middle of an operation (e.g., writing data).
  - **Deferred Cancellation**: Allows the thread to periodically check if it should terminate, ensuring more orderly thread termination.
- **Signal Handling:**
  - Signals in UNIX systems notify processes of events. Asynchronous signals are generated outside the process, while synchronous signals originate from the process itself.
  - Signal handling involves generating, delivering, and handling the signal through a signal handler.

## Thread Pools

A thread pool is a group of pre-created threads that can be used as needed, providing performance benefits in applications where threads are frequently created and destroyed (e.g., servers that spawn a new thread for each client connection).

- **Advantages:**
  - Faster service, as threads are reused rather than created from scratch.
  - Bound the number of threads to prevent resource exhaustion.
- **Disadvantages:** If the thread pool is empty, the system must wait for threads to become available.

## CPU Scheduling

### Scheduling Criteria

- **CPU Utilization:** Keep the CPU busy.
- **Throughput:** Number of processes completed per time unit.
- **Burst Time:** Time consumed by each process.
- **Turnaround Time:** Total time to execute a process (including waiting time).
- **Waiting Time:** Time a process spends in the ready queue.
- **Response Time:** Time from submitting a request until the first response (for time-sharing systems).
- **Arrival Time:** Time when a process arrives in the system for execution.

### Scheduling Algorithm Optimization Criteria

- Maximize CPU utilization.
- Maximize throughput.
- Minimize turnaround time.
- Minimize waiting time.
- Minimize response time.

## CPU Scheduling Algorithms

1. **First- Come, First-Served (FCFS)**
2. **Round Robin**
3. **Shortest Job First (SJF) – Preemptive**
4. **Shortest Job First (SJF) – Non Preemptive**
5. **Priority Scheduling - Preemptive**
6. **Priority Scheduling - Non Preemptive**

## Preemptive & Non-Preemptive Scheduling

### Preemptive Scheduling

- The CPU can be preempted from a running process to handle a higher-priority process.
- Supports multitasking by allowing the system to handle multiple processes at once.
- **Examples:** SRTF (Shortest Remaining Time First), Round Robin, Priority Scheduling.

### Non-Preemptive Scheduling

- Once a process starts, it must finish before another process is executed.
- Cannot handle multitasking as it processes in a sequential manner.
- **Examples:** FCFS (First Come First Serve), SJF (Shortest Job First), Priority Scheduling.

Preemptive Scheduling	Non-Preemptive Scheduling
Processor can be preempted to execute a different process in the middle of execution of any current process.	Once Processor starts to execute a process it must finish it before executing the other. It cannot be paused in middle.
CPU utilization is more compared to Non-Preemptive Scheduling.	CPU utilization is less compared to Preemptive Scheduling.
Waiting time and Response time is less.	Waiting time and Response time is more.
The preemptive scheduling is prioritized. The highest priority process should always be the process that is currently utilized.	When a process enters the state of running, the state of that process is not deleted from the scheduler until it finishes its service time.
If a high priority process frequently arrives in the ready queue, low priority process may starve.	If a process with long burst time is running CPU, then another process with less CPU burst time may starve.
Preemptive scheduling is flexible.	Non-preemptive scheduling is rigid.
Ex:- SRTF, Priority, Round Robin, etc.	Ex:- FCFS, SJF, Priority, etc.

## System Calls Related to Scheduling

System calls allow processes to modify their priorities and scheduling policies. Users can generally lower their process priority but require superuser privileges to increase it or modify other users' processes.

1. **nice()** – Changes the process priority. Negative values increase priority (superuser required), positive values decrease it.
2. **getpriority()** – Returns the highest priority of specified processes.
3. **setpriority()** – Sets the priorities of specified processes. Superuser needed to lower priorities.
4. **sched\_setscheduler()** – Sets scheduling policy and parameters for a process (pid).
5. **sched\_getscheduler()** – Returns the scheduling policy of a process.
6. **sched\_setparam()** – Sets scheduling parameters for a thread.
7. **sched\_getparam()** – Fetches scheduling parameters of a thread.
8. **sched\_get\_priority\_max()** – Returns the maximum priority for a scheduling policy.
9. **sched\_get\_priority\_min()** – Returns the minimum priority for a scheduling policy.
10. **sched\_rr\_get\_interval()** – Fetches quantum used for round-robin scheduling.

## PYQ

**How user level threads are different from kernel level threads? Under what circumstances is one type “better” than the other?**

Aspect	User-Level Threads (ULT)	Kernel-Level Threads (KLT)
Management	Managed by user-level library.	Managed by the OS kernel.
Performance	Faster due to no kernel involvement.	Slower due to kernel intervention.
Concurrency	OS sees a single process, limiting concurrency.	OS schedules threads independently, enabling better concurrency.
Blocking	If one thread blocks, all block.	Other threads can continue executing if one blocks.
Portability	More portable, independent of kernel.	Less portable, requires kernel support.

### Circumstances Where One is Better Than the Other

- **ULT is better** when low overhead and quick thread management are needed (e.g., I/O-bound tasks).
- **KLT is better** for full concurrency and multi-core utilization, such as CPU-bound tasks where threads need independent scheduling.

## Advantages and Disadvantages of Multiple Threads vs Multiple Processes

### Advantages

1. **Efficiency:** Threads share memory, making data sharing faster than separate processes.
2. **Lower Overhead:** Creating and managing threads is less costly than processes.

### Disadvantage

- **Synchronization Issues:** Threads share memory, causing potential data corruption unless properly synchronized.

### Application Example

- **Threads benefit:** Web servers (handle concurrent requests).
- **Threads don't benefit:** Video rendering (sequential processing).

## Resources Used by Threads vs Processes

### Resources Used by Threads

1. **CPU Context:** Each thread has its own set of registers and program counter to keep track of its execution state.
2. **Thread-Specific Data:** This includes the stack, which is used for local variables, function calls, and return addresses.
3. **Shared Resources:** Threads within the same process share resources like memory space and file descriptors, making it more efficient in terms of inter-thread communication.

### Resources Used by Processes

1. **Memory Space:** Each process has its own independent memory space, which is isolated from other processes. This includes the heap and stack used by the process.
2. **Process Control Block (PCB):** The PCB holds all information about the process, such as its state, program counter, CPU registers, and memory management information.
3. **File Descriptors:** Processes use file descriptors to manage access to files or I/O devices. Each process has its own set of file descriptors, which are typically inherited by child processes.

### Differences Between Threads and Processes

- **Memory Space:**
  - **Threads:** Share the same memory space within a process, which allows them to efficiently communicate with each other.
  - **Processes:** Each process has its own separate memory space, ensuring better isolation but requiring more overhead for inter-process communication.
- **Overhead:**
  - **Threads:** Have lower overhead because they share memory and other resources within the same process.
  - **Processes:** Have higher overhead due to the need for separate memory spaces and independent control blocks, as well as the complexity of inter-process communication.
- **Resource Usage:**
  - **Threads:** Consume fewer system resources, as only the thread's context and stack need to be maintained.
  - **Processes:** Require more system resources, as each process must maintain its own memory, process control block, and other resources.

**Describe the actions taken by a kernel to switch context a.) Among threads.b.) Among processes.**

### Context Switch Among Threads

1. **Save Current Thread State:** The kernel saves the current thread's CPU registers, program counter, and stack pointer.
  2. **Update TCB:** The thread's state is updated in the Thread Control Block (TCB).
  3. **Select Next Thread:** The scheduler picks the next thread to run.
  4. **Load New Thread State:** The kernel restores the next thread's registers, program counter, and stack pointer.
  5. **Execute New Thread:** The CPU executes the selected thread.
- **Fast:** Thread context switches are quick as threads share memory and resources.

### Context Switch Among Processes

1. **Save Current Process State:** The kernel saves the process's CPU registers, program counter, and memory mappings.
  2. **Update PCB:** The Process Control Block (PCB) is updated with the process's state.
  3. **Select Next Process:** The scheduler chooses the next process to run.
  4. **Load New Process State:** The kernel restores the new process's registers and memory context.
  5. **Switch Memory Context:** The memory management unit (MMU) updates the address space.
  6. **Execute New Process:** The CPU resumes executing the new process.
- **Slower:** Process context switching involves switching entire memory contexts, which is more time-consuming.

# Process Synchronization and Deadlock

## Synchronization in Concurrent Programming

Synchronization is the process of coordinating the execution of multiple processes or threads to ensure correct and predictable behavior in a concurrent system. In the context of concurrent programming, synchronization is essential to avoid conflicts and ensure that shared resources are accessed in a controlled and orderly manner.

## Race Condition

A **race condition** occurs when multiple processes or threads access shared resources concurrently and the outcome depends on the timing or order of execution. This can lead to unexpected or incorrect behavior, as the processes interfere with each other in an unpredictable manner.

### Key Characteristics:

1. **Shared Resources:** Race conditions usually involve shared resources such as variables, memory, or files.
2. **Concurrency:** The processes or threads execute in parallel, making it difficult to predict their interaction.
3. **Non-deterministic Outcome:** The final result can vary depending on the order of execution, which might change with every run.

## Critical Section Problem

The **Critical Section** problem arises when multiple processes or threads attempt to access shared resources concurrently. The goal is to design a mechanism that ensures only one process or thread can access the shared resource at a time, thus avoiding data inconsistency and race conditions.

### General structure of process Pi

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```

### Algorithm for Process Pi

```
do {  
    while (turn == j);  
    critical section  
    turn = j;  
    remainder section  
} while (true);
```

### Conditions for a Solution:

- **Mutual Exclusion:** Only one process can be in its critical section at any given time.
- **Progress:** If no process is in the critical section, and multiple processes want to enter, the system should allow one of the processes to enter.
- **Bounded Waiting:** A process should not have to wait indefinitely to enter the critical section.

### Peterson's Solution (Two Process Scenario)

- **Variables:**
  - turn: indicates whose turn it is to enter the critical section.
  - flag[]: indicates if a process is ready to enter the critical section.
- **Algorithm:**
  - Each process sets flag[i] to true and turn to the other process. It then waits until the other process is not ready or it's its turn.

## Synchronization Tools

### Mutex Locks

A **Mutex lock** is one of the simplest synchronization tools designed to solve the critical section problem by ensuring that only one process at a time can enter a critical section. Here's how it works:

- **acquire()**: A process must acquire the lock before entering the critical section. It checks whether the lock is available and, if not, waits until it becomes available (often using busy-waiting).
- **release()**: After finishing its task in the critical section, the process releases the lock, making it available for other processes to acquire.

```
acquire() {  
    while (!available); // Busy wait  
    available = false;  
}  
  
release() {  
    available = true;  
}
```

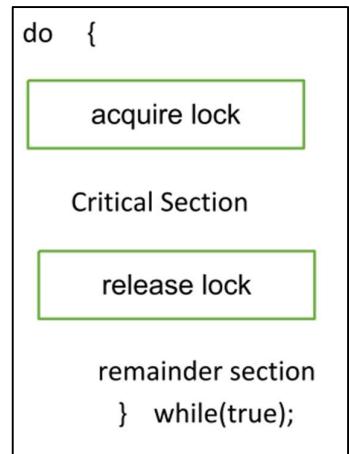
A **spinlock** is a type of mutex lock where the process "spins" (busy waits) while waiting for the lock to become available. The major disadvantage of spinlocks is that they consume CPU cycles while waiting, making them inefficient, especially when the critical section is large or the lock is held for a long time.

### Solution to Critical Section Using Mutex Locks

The general structure for using a **mutex lock** to solve the critical section problem is as follows:

In this solution:

1. **Acquire**: Before entering the critical section, the process checks if the lock is available. If the lock is unavailable (another process is in the critical section), the process waits.
2. **Critical Section**: Once the lock is acquired, the process can safely perform operations in the critical section without interference from other processes.
3. **Release**: After the critical section, the process releases the lock, allowing other processes to enter the critical section.



This solution ensures mutual exclusion, meaning only one process can be in the critical section at any given time.

## Semaphore

A **Semaphore** is a more advanced synchronization tool that can handle more complex scenarios than mutex locks. It is an integer variable used to control access to shared resources. Semaphores support two atomic operations: **wait()** and **signal()**.

### Types of Semaphores:

1. **Counting Semaphore**: Can have any integer value, useful for controlling access to a pool of resources.
2. **Binary Semaphore**: Can only be 0 or 1, similar to a mutex lock.

### Definitions of Semaphore Operations:

- **wait(S)**: This operation decrements the semaphore. If the semaphore value is less than or equal to zero, the process waits (busy waits) until the semaphore is positive again.
- **signal(S)**: This operation increments the semaphore, signaling that a process has finished using a resource and it's now available for others.

```
wait(s) {  
    while (s <= 0)  
        ; // Busy wait  
    s--;  
}
```

```
signal(s) {  
    s++;  
}
```

### Semaphore Usage:

Semaphores can be used to solve synchronization problems, such as coordinating the execution order of processes.

### Example:

Consider two processes: P1 with statement S1 and P2 with statement S2. We require that S2 be executed only after S1 has completed.

- **Solution**: Use a semaphore synch, initialized to 0. After P1 executes S1, it signals the semaphore. P2 waits for the semaphore to be signaled before executing S2.

### Counting Semaphores:

A counting semaphore can be used to manage a resource pool, such as controlling access to a limited number of identical resources. For example, a semaphore initialized to the number of available resources will be decremented when a process acquires a resource and incremented when a process releases it.

## **Classic Synchronization Problems**

Several well-known synchronization problems are used to evaluate different synchronization schemes. Some of the classical problems include:

### **a. Bounded-Buffer Problem (Producer-Consumer Problem)**

In this problem, a producer produces items and stores them in a buffer. A consumer consumes items from the buffer. The challenge is to ensure that the producer and consumer do not overwrite each other's work or attempt to access an empty buffer.

#### **Solution:**

- Use semaphores to manage the buffer's empty and full slots and ensure mutual exclusion for accessing the buffer.

### **b. Readers-Writers Problem**

In this problem, a data set is shared by multiple processes, and some of the processes are **readers** (who only read the data) while others are **writers** (who can both read and write). The problem is to ensure that:

- Multiple readers can read concurrently.
- Only one writer can write at a time, and writers have exclusive access to the data.

#### **Solution:**

- Use semaphores and counters to manage the number of readers and writers and provide mutual exclusion.

### **c. Dining Philosophers Problem**

In this problem, multiple philosophers are sitting at a table with a bowl of rice and a chopstick between each pair of philosophers. To eat, a philosopher needs two chopsticks. The challenge is to ensure no deadlock or starvation occurs while multiple philosophers are trying to eat at the same time.

#### **Solution:**

- Use semaphores to manage the chopsticks and ensure that philosophers pick them up in an order that prevents deadlock.

## 1. Bounded-Buffer Problem (Producer-Consumer)

In this problem, we have **n buffers**, where each buffer can hold one item. The goal is to synchronize the **producer** and **consumer** processes while ensuring that no buffer is overfilled or underfilled.

### Semaphore Initialization:

- **mutex**: A binary semaphore initialized to 1 (for mutual exclusion).
- **full**: A counting semaphore initialized to 0 (indicating how many slots are filled).
- **empty**: A counting semaphore initialized to n (representing how many slots are empty).

#### Producer Process:

The producer produces items and places them into the buffer, ensuring synchronization using the semaphores.

```
c

do {
    ...
    /* produce an item in next_produced */
    ...

    wait(empty);    // wait until there are empty slots
    wait(mutex);    // ensure mutual exclusion
    ...

    /* add next produced item to the buffer */
    ...

    signal(mutex); // release mutual exclusion
    signal(full);  // indicate that a slot has been filled
} while (true);
```

#### Consumer Process:

The consumer removes items from the buffer and consumes them.

```
c

do {
    wait(full);    // wait until there is something to consume
    wait(mutex);    // ensure mutual exclusion
    ...

    /* remove an item from buffer to next_consumed */
    ...

    signal(mutex); // release mutual exclusion
    signal(empty); // indicate that a slot has been emptied
    ...

    /* consume the item in next_consumed */
    ...

} while (true);
```

## 2. Readers-Writers Problem

In the Readers-Writers problem, multiple **readers** can access shared data concurrently, but **writers** require exclusive access. The goal is to ensure:

- Multiple readers can read data at the same time.
- Only one writer can access data at a time.
- Writers should have priority when they are ready.

### Semaphore Initialization:

- `rw_mutex`: A semaphore initialized to 1 (for mutual exclusion of writers).
- `mutex`: A semaphore initialized to 1 (for managing the count of readers).
- `read_count`: An integer initialized to 0 (keeps track of how many readers are currently reading).

#### Writer Process:

A writer acquires the `rw_mutex` to ensure exclusive access to the data.

```
c
Copy code

do {
    wait(rw_mutex);
    ...
    /* writing is performed */
    ...
    signal(rw_mutex);
} while (true);
```

#### Reader Process:

A reader increments the `read_count` and acquires the `rw_mutex` only when it's the first reader. It releases the lock when the last reader is done.

```
c
Copy code

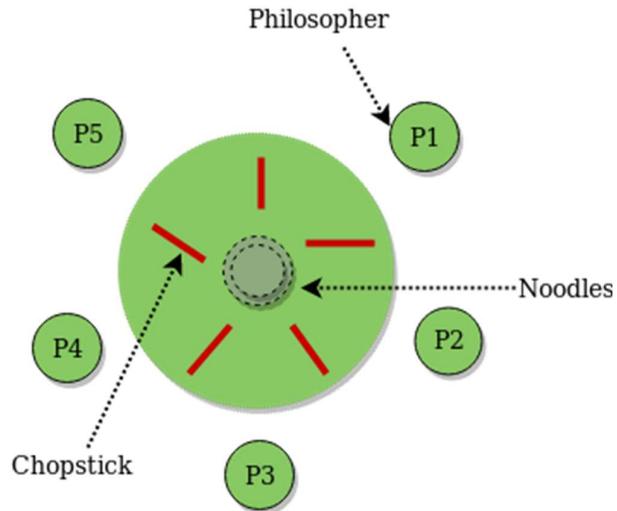
do {
    wait(mutex);
    read_count++;
    if (read_count == 1)
        wait(rw_mutex);
    signal(mutex);

    ...
    /* reading is performed */
    ...

    wait(mutex);
    read_count--;
    if (read_count == 0)
        signal(rw_mutex);
    signal(mutex);
} while (true);
```

### 3. Dining Philosophers Problem

In this problem, philosophers alternate between thinking and eating. To eat, a philosopher needs two chopsticks. The challenge is to synchronize their actions to avoid **deadlock** and **starvation**.



#### Semaphore Initialization:

- `chopstick[5]`: An array of semaphores initialized to 1 (representing 5 chopsticks).

#### Philosopher Process:

Each philosopher picks up two chopsticks (one at a time), eats, and then releases them.

```
c

do {
    wait(chopstick[i]);
    wait(chopstick[(i + 1) % 5]);
    // eat
    signal(chopstick[i]);
    signal(chopstick[(i + 1) % 5]);
    // think
} while (TRUE);
```

[Copy code](#)

## Deadlock and Starvation

### Deadlock

**Deadlock** occurs when two or more processes are waiting indefinitely for each other to release resources. This creates a cyclic dependency and halts all involved processes.

- **Conditions for Deadlock:**

1. Mutual Exclusion: Only one process can use a resource at a time.
2. Hold and Wait: A process holding one resource is waiting for another.
3. No Preemption: Resources cannot be forcibly taken from processes.
4. Circular Wait: A set of processes are waiting for each other in a circular chain.

### Starvation

**Starvation** occurs when a process is indefinitely delayed in obtaining resources because other processes are consistently given priority. This often happens in priority-based scheduling systems when a low-priority process is repeatedly preempted by higher-priority processes.

- **Solution:** Implement **fairness policies** such as **aging** to prevent starvation.

## Methods for Handling Deadlocks

- **Deadlock Prevention:**

- **Mutual Exclusion:** Prevent if possible, but often necessary for certain resources.
- **Hold and Wait:** Require processes to request all resources upfront.
- **No Preemption:** Allow resources to be preempted if necessary.
- **Circular Wait:** Impose a strict order on resource allocation.

- **Deadlock Avoidance:**

- Requires knowledge of the maximum resources each process might need.
- **Banker's Algorithm:** Checks if granting a request would leave the system in a safe state, avoiding potential deadlocks.
- Ensures the system is in a **safe state**, where processes can finish without causing deadlocks.

## Resource-Allocation Graph (RAG)

A **Resource-Allocation Graph (RAG)** is a graphical representation of the allocation of resources to processes and can be used to detect deadlock.

- **Nodes** in the graph represent:
  - **Processes** (denoted as circles).
  - **Resources** (denoted as squares).
- **Edges** in the graph represent:
  - **Request edges** (represented by a directed edge from a process to a resource): Indicates that a process is requesting a resource.
  - **Assignment edges** (represented by a directed edge from a resource to a process): Indicates that a resource has been allocated to a process.

### Conditions for Deadlock in a RAG:

- **Deadlock occurs** if there is a cycle in the graph, where each process is waiting for a resource held by the next process in the cycle. This cycle forms a circular wait, one of the necessary conditions for deadlock.

### Example of a Resource-Allocation Graph:

1. Process **P1** requests resource **R1** (represented by an edge from P1 to R1).
2. Process **P2** holds resource **R1** and is requesting resource **R2** (represented by an edge from P2 to R2).
3. Process **P3** holds resource **R2** and is requesting resource **R1** (represented by an edge from P3 to R1).

In this case, there is a cycle:  $P1 \rightarrow R1 \rightarrow P2 \rightarrow R2 \rightarrow P3 \rightarrow R1$ . This cycle indicates a **deadlock** because each process is waiting for a resource held by another, forming the circular wait condition.

## Safe and Unsafe States

- **Safe state:** A state where there is at least one safe sequence (a sequence of process executions where each process can finish its execution without causing a deadlock).
- **Unsafe state:** A state where no such sequence exists, and granting further resource requests might lead to deadlock.

## Recovery from Deadlock

### 1. Process Termination

- **Abort all deadlocked processes:** Terminate all processes involved in the deadlock.
- **Abort one process at a time:** Terminate processes one by one until the deadlock is resolved.

### 2. Resource Preemption

- **Selecting a Victim:** Preempt resources from a selected process to break the deadlock.
- **Rollback:** Revert a process to a previous safe state and restart it.
- **Starvation Prevention:** Ensure no process is repeatedly chosen as a victim to prevent starvation.

## Criteria for Selecting a Victim Process

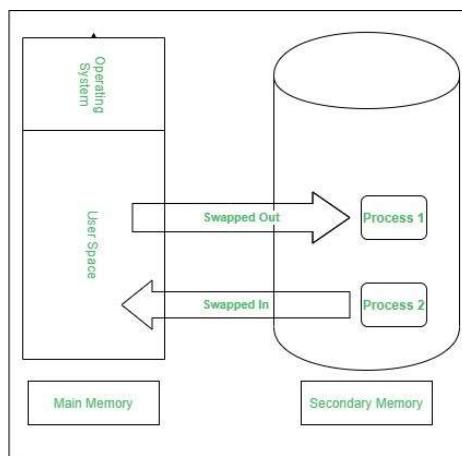
- **Priority:** Abort lower-priority processes.
- **Time Computed:** Terminate processes that have not made much progress.
- **Resource Usage:** Abort processes using fewer resources.
- **Remaining Resource Needs:** Favor processes that are close to completion.
- **Number of Processes Terminated:** Minimize the number of terminated processes.
- **Interactive vs. Batch:** Prioritize interactive processes to reduce user impact.

# Memory Management

## Swapping

Swapping involves temporarily moving a process out of memory to a backing store (usually a disk) and bringing it back when needed, allowing more processes to run than the physical memory can hold.

- **Backing Store:** A large, fast disk that can store memory images of all processes, providing quick access when processes need to be swapped in or out.
- **Roll-out, Roll-in:** A variant of swapping for priority-based scheduling. Lower-priority processes are swapped out to make room for higher-priority ones.
- **Swap Time:** The time it takes to swap a process is proportional to the amount of memory being transferred. A significant portion of swap time comes from the transfer itself.
- **Context Switch Time with Swapping:** When a process needs to be swapped in or out, the context switch time increases. For example, swapping a 100MB process with a transfer rate of 50MB/sec would take 4 seconds (2 seconds for swap out + 2 seconds for swap in).



### Advantages:

- Allows execution of multiple processes despite limited memory.
- Utilizes main memory effectively.
- Enables the concept of virtual memory and optimizes scheduling.

### Disadvantages:

- Risk of data loss if power is lost during swapping.
- Can lead to page faults and reduced performance.

## Contiguous Memory Management

Contiguous Memory Management is a technique where processes are allocated memory in a continuous, sequential manner. The main idea is to assign a single block of memory to each process based on its size.

### Types of Contiguous Memory Management:

#### 1. Fixed (Static) Partitioning:

Memory is divided into a fixed number of partitions, each capable of holding one process. The number of partitions and their size are fixed, meaning the degree of multi-programming is limited by the number of partitions.

- **Advantages:** Simple allocation and management.
- **Disadvantages:**
  - Internal fragmentation: If a process is smaller than the partition, the unused memory within that partition is wasted.
  - Restricted maximum process size: A process cannot exceed the size of the partition.
  - Degree of multi-programming is restricted by the number of partitions.

#### 2. Variable (Dynamic) Partitioning:

Initially, memory is a single free block. When processes arrive, the memory is divided into partitions according to the process size. This allows for more flexibility as processes can vary in size.

- **Advantages:** No internal fragmentation (as partition size matches the process size).
- **Disadvantages:**
  - External fragmentation: Small gaps of free memory between allocated blocks can occur.
  - **Solution:** Compaction (moving processes to make free memory contiguous), but this is resource-intensive.

## **Memory Fragmentation:**

- **Internal Fragmentation:** Wasted space inside a partition in the Fixed Partitioning scheme.
- **External Fragmentation:** Wasted space between partitions in the Variable Partitioning scheme.

## **Solutions to Fragmentation:**

1. **Compaction:** Rearranging memory to create a large contiguous block of free memory. However, it requires moving processes and can cause performance overhead.
2. **Non-contiguous Memory Allocation:** Dividing memory into fixed-size blocks called frames and processes into pages. This can eliminate external fragmentation.

## **Advantages of Contiguous Memory Management:**

- Simple to monitor available memory blocks.
- Good read performance as files can be read in a single session.
- Easy setup and management.

## **Disadvantages of Contiguous Memory Management:**

- Fragmentation issues (especially with Variable Partitioning).
- Need to know the final size of a process to select the right partition.
- Extra space in holes must be managed or used efficiently.

## Paging

Paging is a memory management scheme that eliminates the issues of external fragmentation and large chunks of varying-sized memory allocations. It allows for the non-contiguous allocation of physical memory to processes, providing greater flexibility and efficiency.

### Key Concepts:

- **Physical Memory:** Divided into fixed-sized blocks called **frames**. The size of a frame is typically a power of 2 and varies between 512 bytes to 16 MB.
- **Logical Memory:** Divided into blocks of the same size called **pages**. The size of a page is the same as the frame size and is defined by the hardware.

### Address Translation:

- **Page Number (p):** The part of the logical address that acts as an index into the **page table**. The page table contains the base address of each page in physical memory.
- **Page Offset (d):** The part of the logical address used to find the specific location within the page.

Given:

- **Logical Address Space:**  $2^{m-n}$
- **Page Size:**  $2^n$

The address generated by the CPU is divided into two parts:

- The higher-order  $m-n$  bits of the address designate the **page number**.
- The lower-order  $n$  bits designate the **page offset**.

The page number is used as an index into the **page table**, and the base address from the page table is combined with the page offset to form the **physical address**.

### Paging Hardware:

- The **page table** holds the **base address** of each page in physical memory.
- When a page is accessed, the page number is used to index into the page table to obtain the base address of the corresponding frame. The **page offset** is then added to this base address to get the physical address.

### Paging Example:

- **Page Size:** 4 bytes
- **Physical Memory:** 32 bytes (8 frames)
- **Process Size:** 16 bytes (4 pages)

For logical address 0 (binary: 0000), it refers to **Page 0**, and the offset is 0. If Page 0 is mapped to Frame 5, the physical address is: Physical Address =  $(5 \times 4) + 0 = 20$

For logical address 3 (binary: 0011), it is still in **Page 0**, but with an offset of 3:

Physical Address =  $(5 \times 4) + 3 = 23$

For logical address 4 (binary: 0100), it maps to **Page 1**, and according to the page table, Page 1 is mapped to Frame 6: Physical Address =  $(6 \times 4) + 0 = 24$

### Demand Paging

- **Pure Demand Paging:** In this scheme, a process begins with no pages in memory. When the process tries to access a page that is not in memory, a **page fault** occurs. The operating system will then load the required page into memory from the secondary storage (disk).
- Demand paging works well with the **locality of reference**, which means that programs tend to access a small set of memory locations repeatedly, thus reducing the number of page faults.

### Hardware Support for Demand Paging:

- **Page Table:** Stores the **valid/invalid bit** that indicates whether a page is in memory or not.
- **Secondary Memory:** Uses **swap space** for storing pages that are not in use.
- **Instruction Restart:** When a page fault occurs, the CPU can restart the instruction once the page is loaded into memory.

### What Happens if There is No Free Frame?

- Physical memory could be full, and no free frames are available for new pages. In such a case, a **page replacement algorithm** is needed.
- The operating system has to decide which page to **swap out** (i.e., move to secondary memory) to make room for the new page.

## Page Replacement Algorithm

The objective is to minimize the number of page faults. Some common algorithms include:

1. **FIFO (First In First Out)**: The oldest page is swapped out first.
2. **LRU (Least Recently Used)**: The page that has not been used for the longest time is swapped out.
3. **Optimal**: The page that will not be needed for the longest time in the future is swapped out.

## Swap Space

**Swap Space** is a portion of disk storage used by the operating system to extend physical RAM. When RAM is full, inactive data is moved to the swap space to make room for active processes. This allows the system to handle more processes than can fit in physical memory.

### Types of Swap Space:

4. **Swap Partition**:
  - A dedicated portion of the disk set aside specifically for swapping data. It is typically faster than a swap file because it is directly allocated and managed by the operating system.
4. **Swap File**:
  - A file created on the existing file system that serves as swap space. This is more flexible because it can be resized easily, but may have slightly slower performance than a dedicated partition.

### How Swap Space Works:

- **When RAM is full**: If there is more demand for memory than physical RAM can handle, the operating system will start swapping out less-used pages from RAM to the swap space.
- **When a page is needed**: If a process needs data that has been swapped out to the disk, a page fault occurs. The operating system swaps in the required page from the swap space, and another page might be swapped out to make room.

## PYQ

**Explain the following terms in brief 1)Demand Paging 2)LRU 3)Belady's Anomaly**

### **1. Demand Paging**

Demand Paging is a memory management technique where pages are only loaded into physical memory when they are required, typically due to a page fault. Rather than loading the entire program into memory at once, only the necessary pages are loaded, which reduces memory usage and allows more processes to run concurrently. When a page that is not in memory is accessed, a page fault occurs, and the operating system retrieves the page from secondary storage.

**Advantages:**

- Efficient memory utilization.
- Faster program startup.

**Disadvantages:**

- Page faults introduce delays.
- Fragmentation can occur.

### **2. LRU (Least Recently Used)**

LRU is a page replacement algorithm that swaps out the page that has not been accessed for the longest period when memory is full. It operates on the principle that recently used pages are more likely to be used again, while older pages are less likely to be needed.

**Implementation:**

- **Queue-based:** Move the most recently accessed page to the front.
- **Timestamp-based:** Track the last access time of each page.

**Advantages:**

- Intuitive and effective in many cases.
- Often performs well in systems with predictable access patterns.

**Disadvantages:**

- High overhead to track access.
- May not be optimal in all scenarios.

### **3. Belady's Anomaly**

Belady's Anomaly refers to a phenomenon where increasing the number of page frames can unexpectedly increase the number of page faults, especially in the FIFO (First-In, First-Out) page replacement algorithm. This counterintuitive behavior occurs because FIFO doesn't consider the recency or frequency of page accesses, leading to inefficient page replacements.

#### **Example:**

- Increasing frames from 3 to 4 might cause more page faults in some cases when using FIFO.

#### **Significance:**

- Shows that increasing memory does not always improve performance.
- Emphasizes the importance of choosing an appropriate page replacement algorithm.

4. **Virtual Memory:** Virtual memory is a memory management technique that allows the operating system to use a combination of physical RAM and disk space to create an "illusion" of a larger amount of memory than what is physically available. It enables programs to run even if there is not enough physical memory by swapping data between RAM and disk storage (usually a swap file or paging file).
5. **Compaction:** Compaction is the process of rearranging the memory in a system to eliminate fragmentation. In systems that use dynamic memory allocation, free memory blocks can become scattered, leading to inefficient use of memory. Compaction moves data to combine free memory spaces into a larger contiguous block, allowing for better memory utilization.
6. **Thrashing:** Thrashing occurs when a computer's virtual memory system becomes overwhelmed by excessive paging or swapping between RAM and disk. This happens when there isn't enough physical memory to handle the active processes, causing the system to spend more time swapping data in and out of memory than executing processes, leading to a significant slowdown.

# File Management & Access Control

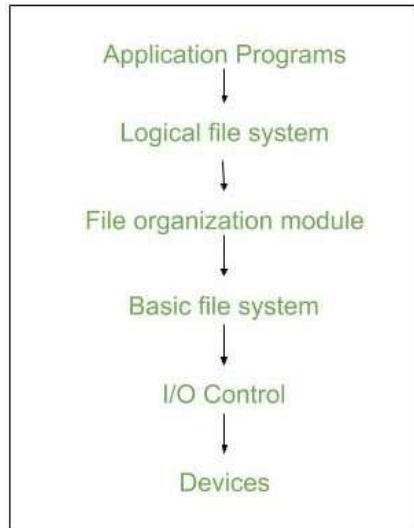
## File Attributes

- **Name:** Human-readable identifier.
- **Identifier:** Unique numeric tag used by the system to identify the file.
- **Type:** Indicates file format (e.g., text, binary) in systems that support types.
- **Location:** Pointer to the file's location on storage.
- **Size:** Current size of the file in bytes.
- **Protection:** Permissions that define who can read, write, or execute the file.
- **Time, Date, and User Identification:** Metadata for security, protection, and usage tracking (e.g., creation date, last modified date).
- **Extended Attributes:**
  - **Checksum:** Ensures file integrity.
  - Other custom attributes based on system requirements.

## File Operations

1. **Create:** Allocates space and metadata for a new file.
2. **Write:** Adds data at the current write pointer.
3. **Read:** Retrieves data from the current read pointer.
4. **Reposition (Seek):** Moves the file pointer to a specific position (does not involve I/O).
5. **Delete:** Removes the file from the directory structure.
6. **Truncate:** Clears the file's contents but retains its attributes and structure.
7. **Open:**
  - Searches the directory for the file.
  - Loads its metadata into memory for access.
8. **Close:**
  - Writes any updated metadata back to disk.
  - Removes the file's in-memory data.

## Layers in File System



### 1. Application Programs:

- User interface layer for file operations (e.g., create, delete, read).
- Examples: Text editors, file browsers, command-line tools.

### 2. Logical File System:

- Manages metadata through File Control Blocks (FCBs).
- Tracks owner, size, permissions, and file location details.

### 3. File Organization Module:

- Maps logical blocks to physical blocks.
- Tracks free space for allocation.

### 4. Basic File System:

- Handles read/write commands to physical blocks.
- Manages memory buffers and caches for performance.

### 5. I/O Control Level:

- Device drivers as an interface between OS and hardware.
- Translates block numbers into hardware-specific commands.

### 6. Device Layer:

- Actual hardware performing read/write operations.
- Includes storage devices like HDDs, SSDs, and optical disks.

## File-System Implementation

### 1. On-Disk Structures

- **Boot Control Block:**  
Contains information needed to boot the OS from the volume.  
*Location:* Usually the first block of the volume.
- **Volume Control Block (e.g., Superblock):**  
Stores details like total blocks, free blocks, block size, and free block pointers.
- **Directory Structure:**  
Organizes files with their names and inode numbers (or master file table).
- **File Control Block (FCB):**  
Contains file details like:
  - Inode number
  - Permissions
  - File size
  - Timestamps

### 2. In-Memory Structures

- **Mount Table:**  
Tracks information about mounted volumes.
- **In-Memory Directory Structure:**  
Stores recently accessed directory information for faster lookups.
- **System-Wide Open-File Table:**  
Contains:
  - FCB of open files.
  - Tracks the number of processes accessing each file.
- **Per-Process Open-File Table:**  
Points to entries in the system-wide table and contains process-specific file information.
- **Buffers:**  
Temporarily hold file-system blocks for read/write operations.

## File Operations

### Creating a New File

#### 1. Logical File System:

- Called by the application program to create a file.
- Allocates a new FCB.

#### 2. Directory Update:

- Reads the appropriate directory into memory.
- Updates it with the new file name and FCB.
- Writes back the updated directory to the disk.

### Reading a File

#### 1. open() System Call:

- Checks the **System-Wide Open-File Table**:
  - If file is already in use, creates an entry in the **Per-Process Open-File Table** pointing to it.
  - If not in use, searches the **directory structure** for the file.

#### 2. Caching for Speed:

- Frequently accessed parts of the directory structure are cached.

#### 3. File Control Block (FCB):

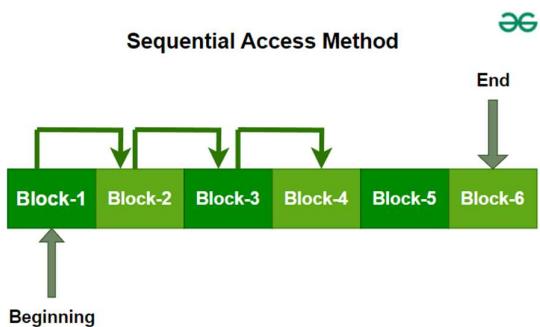
- FCB is loaded into the **System-Wide Open-File Table**.
- Tracks processes using the file.

## File Access Methods in Operating Systems

File access methods determine how data is organized, retrieved, and modified in a file system. The choice of access method can significantly impact system performance and efficiency. There are three primary file access methods: **Sequential Access**, **Direct Access**, and **Index Sequential Access**.

### 1. Sequential Access

In sequential access, data is processed in order, one record after another. The file pointer moves to the next position after each read or write operation. It is simple and commonly used in tasks like text editing or processing large files in a specific order.



- **Advantages:**

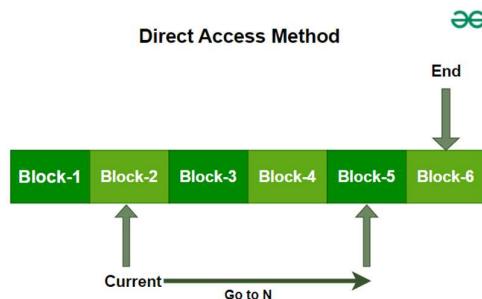
- Simple and easy to implement.
- Ideal for batch processing or tasks that need to access records in sequence.
- Less prone to data corruption since data is written sequentially.

- **Disadvantages:**

- Slow when accessing specific records, as the entire file must be read in sequence.
- Inefficient for frequent updates, as inserting or deleting records requires shifting data.

## 2. Direct Access (Relative Access)

Direct access allows random access to any record in a file. The file is viewed as a sequence of fixed-length blocks, and records can be accessed in any order by providing the block number. This method is often used in databases.



- **Advantages:**

- Faster access time, as records can be directly accessed without reading the previous ones.
- Suitable for applications that require random access to records.

- **Disadvantages:**

- More complex to implement due to the need for sophisticated algorithms.
- Higher storage overhead to maintain the mapping of blocks or records.

## 3. Index Sequential Access

Index Sequential Access combines the benefits of sequential and direct access. It uses an index to store pointers to records in the file, allowing both sequential and random access. The file is processed sequentially, but the index allows for efficient searching.

- **Advantages:**

- Efficient searching through the index.
- Allows both sequential and random access, making it versatile for various applications.
- Improved data management, especially for large files.

- **Disadvantages:**

- Complex to implement and maintain.
- Requires additional storage for the index.
- Updates to the data require updates to the index as well, leading to overhead.

## File Allocation Methods

File allocation methods define how files are stored in disk blocks, ensuring efficient disk space utilization and quick access. The three main methods are **Contiguous Allocation**, **Linked Allocation**, and **Indexed Allocation**, each with distinct advantages and limitations.

### 1. Contiguous Allocation

In contiguous allocation, each file is stored in a contiguous set of blocks. The starting block and the length of the file are specified in the directory entry. For example, if a file starts at block 19 and requires 6 blocks, it will occupy blocks 19 through 24.

#### Advantages:

- Supports both **sequential** and **direct access**, making it fast for both types of access.
- **Minimal seeks** are needed, improving access speed due to the contiguity of blocks.

#### Disadvantages:

- **External fragmentation** occurs when free blocks are scattered, leading to wasted space.
- **Internal fragmentation** happens if the file's allocated space exceeds its actual size.
- Difficult to expand file size as it requires contiguous free space.

### 2. Linked Allocation

Here, each file is stored as a linked list of blocks. The blocks may be scattered throughout the disk, and each block contains a pointer to the next block. The directory entry stores the pointers to the first and last block.

#### Advantages:

- **Flexible file size**: Files can grow easily without needing contiguous space.
- **No external fragmentation** as blocks can be placed anywhere on the disk.

#### Disadvantages:

- **Slower access**: Random access is not possible, requiring sequential traversal of blocks.
- **Pointer overhead**: Each block must store a pointer to the next block, adding extra storage overhead.

### **3. Indexed Allocation**

In indexed allocation, an index block contains pointers to all the blocks used by a file. The directory entry holds the address of the index block, and the index block stores the addresses of all data blocks.

#### **Advantages:**

- Supports **direct access** to any block, which is faster than linked allocation.
- **No external fragmentation** as file blocks can be scattered.

#### **Disadvantages:**

- **Pointer overhead:** Requires a separate index block, which adds additional storage cost.
- **Inefficient for small files:** Small files may waste space in the index block if not enough blocks are used.

## Unix File System (UFS)

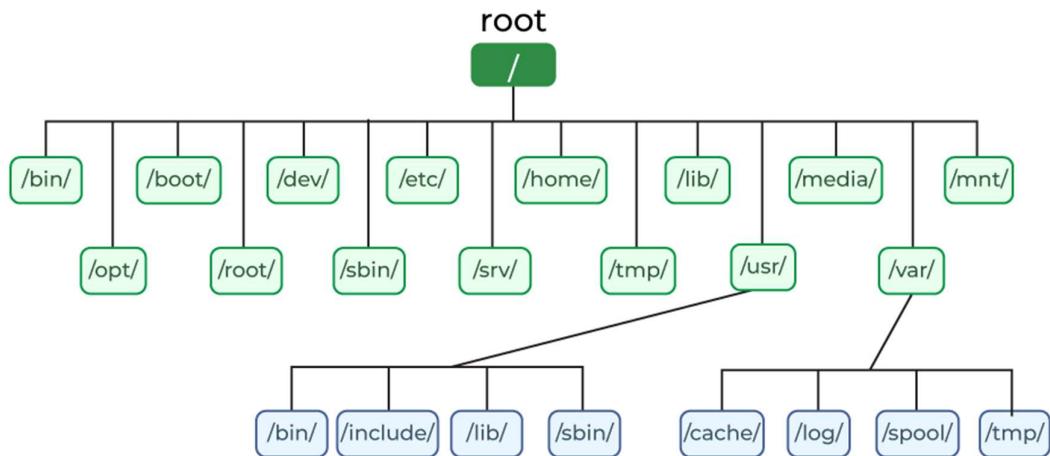
The Unix File System (UFS) is a standard file system used by Unix and Unix-like operating systems. It is designed to manage data in a way that is both efficient and flexible, providing a hierarchical structure for organizing and storing files. Below is an overview of the Unix File System, its structure, and key components.

### Hierarchical Structure

The Unix File System is organized in a tree-like structure, where the topmost directory is called the **root directory (/)**. All other files and directories are stored beneath this root directory, forming a hierarchy.

### Key Components of the Unix File System

- **Root Directory (/):** This is the starting point of the Unix File System. All files and directories reside under this root directory.
- **Directories:** In Unix, directories are special types of files that contain references to other files or directories. Directories organize files and help in navigating the file system.
- **Files:** A file can be an ordinary file containing data, a directory file, or a special file (like a device file). The Unix system treats everything as a file, including hardware devices.



## Types of Files in Unix

Unix recognizes different types of files:

- **Ordinary Files:** Regular files that contain data, such as text files or executable programs.
- **Directories:** Special files that organize other files and directories.
- **Special Files:** Represent hardware devices and I/O operations. These are of two types:
  - **Block Special Files:** Represent devices that perform data transfer in blocks (e.g., disk drives).
  - **Character Special Files:** Represent devices that perform data transfer one character at a time (e.g., keyboards).
- **Pipes:** A method for linking commands together where the output of one command becomes the input of another.
- **Sockets:** Special files used for inter-process communication, often used in client-server applications.
- **Symbolic Links:** A special file that references another file or directory, acting as a shortcut.

## Inodes and File Names

- **Inode:** Each file in Unix is associated with an inode, which contains metadata about the file (e.g., size, owner, permissions, and pointers to the file's data blocks). Inodes are used to track the physical location of files in the file system.
- **File Pathname:** A file's path is a string that describes its location within the file system hierarchy. It can be an absolute pathname (starting from the root /) or a relative pathname (relative to the current directory).

## File Permissions and Security

Unix uses file permissions to control access to files. Each file or directory has permissions for three types of users:

- **Owner:** The user who owns the file.
- **Group:** Users who are in the same group as the file's owner.
- **Others:** All other users.

Permissions include:

- **Read (r)**: Allows viewing the contents of a file.
- **Write (w)**: Allows modifying the contents of a file.
- **Execute (x)**: Allows executing the file as a program.

The permissions are set using the chmod command and can be viewed using the ls -l command.

## File System Operations

Unix provides a set of system calls for manipulating files and directories, such as:

- **open()**: Opens a file.
- **read()**: Reads data from a file.
- **write()**: Writes data to a file.
- **close()**: Closes a file.
- **mkdir()**: Creates a new directory.
- **rmdir()**: Removes a directory.
- **unlink()**: Deletes a file.

## Advantages of the Unix File System

- **Hierarchical Organization**: The tree structure makes it easy to navigate and organize files.
- **Security**: Unix's permission system ensures controlled access to files.
- **Stability**: The file system is known for its robustness and reliability.
- **Flexibility**: Support for symbolic links and mounting external file systems adds flexibility.

## Disadvantages of the Unix File System

- **Complexity**: The Unix File System can be complex, especially for beginners.
- **Command-line Interface**: File system operations are primarily managed through the command line, which may not be as user-friendly as graphical interfaces.

## PYQ

### Role-Based Access Control (RBAC) in Unix

Role-Based Access Control (RBAC) is a method used in Unix and other operating systems to manage and restrict system access based on the roles of individual users within an organization. With RBAC, users are assigned roles that define what operations they can perform on resources, rather than assigning permissions to each user individually. This simplifies administration and enhances security by ensuring that users only have the necessary access for their job functions.

#### Key Concepts:

- **Role:** A set of permissions associated with a specific job function.
- **User:** A person or process that interacts with the system.
- **Permission:** The rights or privileges to access or modify resources (files, directories)

#### How RBAC Works:

In RBAC, each user is assigned a specific role, and each role has defined permissions for resources like files, devices, or directories. The system enforces access control by checking the user's role and associated permissions before allowing access to the requested resource.

#### Example:

Let's consider an organization with the following roles:

- **Admin:** Can perform all operations on all files.
- **Manager:** Can read, write, and execute files in the “/projects” directory.
- **Employee:** Can only read files in the “/projects” directory.

## Explain the different types of files with its usages and examples.

File Type	Description	Examples
Regular Files	Store user data (text, images, binaries)	.txt , .jpg , .exe
Directory Files	Store a list of file names and metadata	/home/ , /etc/
Character Special Files	Represent devices that transfer data one character at a time	/dev/tty , /dev/null
Block Special Files	Represent devices that transfer data in blocks	/dev/sda , /dev/sdb
FIFO (Named Pipes)	Used for inter-process communication	/tmp/myfifo
Socket Files	Used for network communication between processes	/var/run/socket
Symbolic Links	Pointers to other files or directories	/usr/local/bin/somefile
Hard Links	Additional references to the same file (same inode)	ln file1 file2

## Differentiate between Mandatory Locks and Advisory Locks

Feature	Mandatory Locks	Advisory Locks
Enforcement	Enforced by the OS	Voluntary, not enforced by the OS
Process Behavior	Other processes are blocked	Other processes may ignore the lock
Example	Database updates	Text file editing
Use Case	Critical system operations (e.g., DB)	Non-critical or collaborative work